

# **Delegation for Multi-Threaded and Rack-Scale Software**

BY

Noaman Ahmad  
BS CS, Lahore University of Management Sciences, 2018  
MS CS, University of Illinois at Chicago, 2025

## **DISSERTATION**

Submitted as partial fulfillment of the requirements  
for the degree of Doctor of Philosophy in Computer Science  
in the Graduate College of the  
University of Illinois Chicago, 2026

Chicago, Illinois

Defense Committee:  
Jakob Eriksson, Chair and Advisor  
Michael Papka  
Ajay Kshemkalyani  
Luís Pina  
Erdem Koyuncu, Electrical and Computer Engineering  
Xingbo Wu, Microsoft Research

Copyright by

Noaman Ahmad

2026

## **Dedication**

Dedicated to my family, who always have my back.

## Acknowledgements

I am deeply grateful to many individuals who have helped me along this journey.

First, I would like to thank my advisor, Dr. Jakob Eriksson, for his clear guidance, remarkable patience and unwavering support. He taught me how to think like a researcher and approach problems in computer systems. He always made me feel smart during our conversations and encouraged me to follow my heart. Without him, this dissertation would not be possible.

I would also like to thank Dr. Brent Stephens for his help and insights during my first two years at UIC. His presence during our weekly meetings during that time had been very reassuring, and I'll always appreciate that early support.

In addition, I would like to extend my gratitude to the member of my defense committee. Dr. Erdem Koyuncu, Dr. Luís Pina, Dr. Ajay Kshemkalyani, Dr. Michael Papka, and Dr. Xingbo Wu have all been very helpful and understanding. Their diverse perspective has been invaluable for this work.

I am extremely grateful to Dr. Junaid Haroon Siddiqui and Dr. Fareed Zaffar for their guidance during undergrad and their encouragement to pursue higher studies.

My thanks also go to the members of my research group including Ben Baenen, Dr. Nilanjana Basu, Dr. Chen Chen, and Jacob Cohen for the many conversations, both intellectual and casual, that helped me grow as a person and a scientist.

I am thankful to Ahsen Ghafoor, Taha Moaz, Dr. Ahmed Akhtar, Umar pervaiz, Danyal Saeed, and Dr. Usama Muneeb, for their friendship, the casual conversations, the jokes and also the help whenever I needed them.

Finally, I am grateful for the love and support of my family, who have helped me at every step of my journey. From the sacrifices of my parents to the emotional support of my siblings, and to my aunt and uncle who welcomed me into their home every year since I arrived in the U.S., I cannot express how much it all means to me. I will cherish it for the rest of my life.

NA

## **Contributions of Authors**

I would like to acknowledge Dr. Jakob Eriksson and Ben Baenen for their integral role in the completion of this dissertation.

Chapter 3 includes work done in collaboration with Dr. Jakob Eriksson, Ben Baenen and Dr. Chen Chen, and comes from a preprint paper in which I am the first author [7]. The design and implementation of *Trust<T>* is a result of Dr. Eriksson's, Ben's, and my work together and Dr. Chen helped with experimentation and writing.

Chapter 4 builds on work from Chapter 3. Dr. Eriksson helped with its writing and Dr. Chen helped with evaluating the system.

## Contents

Chapter 1	Introduction . . . . .	1
1	<i>Trust</i> < <i>T</i> > . . . . .	2
2	<i>Gossamer</i> . . . . .	4
3	Dissertation Organization . . . . .	6
Chapter 2	State of the Art . . . . .	8
1	Synchronizing shared memory systems . . . . .	8
2	Distributed Systems . . . . .	11
Chapter 3	<i>Trust</i> < <i>T</i> > . . . . .	14
1	Summary . . . . .	14
2	Introduction . . . . .	14
3	<i>Trust</i> < <i>T</i> >: The Basics . . . . .	16
3.1	Trust: a reference to an object . . . . .	16
3.2	Trustee - a thread in charge of entrusted properties . . . . .	17
3.3	Fiber - a delegation-aware, light-weight user thread . . . . .	17
3.4	Delegated context . . . . .	18
4	Core API . . . . .	18
4.1	<code>apply()</code> : synchronous delegation . . . . .	18
4.2	<code>apply_then()</code> : non-blocking delegation . . . . .	19
4.3	<code>launch()</code> : apply in a trustee-side fiber . . . . .	20
5	Key Design and Implementation Details . . . . .	25
5.1	Delegating Closures . . . . .	25
5.2	Scheduling Delegation Work . . . . .	26
5.3	Request and Response Slot Structure . . . . .	27
5.4	Reference counting for Trusts . . . . .	29
6	Evaluation . . . . .	30
6.1	Fetch and Add: Throughput . . . . .	31
6.2	Fetch and Add: Latency . . . . .	35
6.3	Concurrent key-value store . . . . .	36
7	Legacy Application: Memcached . . . . .	41
7.1	Evaluation . . . . .	44
8	Hybrid Delegation . . . . .	48
9	Conclusion . . . . .	48
Chapter 4	<i>Gossamer</i> . . . . .	49
1	Summary . . . . .	49
2	Introduction . . . . .	50

3	Background on RDMA . . . . .	51
3.1	RDMA API . . . . .	52
3.2	Transport types . . . . .	53
4	Gossamer Design . . . . .	54
4.1	Gossamer Daemon . . . . .	55
4.2	Gossamer Process . . . . .	56
4.3	Trustee, Trust and Property . . . . .	56
4.4	Pointers and Serialized Data . . . . .	57
4.5	Request size and TCP fallback . . . . .	59
4.6	Request/Response Slot Format . . . . .	60
4.7	Trustee and Client memory layout . . . . .	63
4.8	Scaling impact of increasing number of Queue Pairs . . . . .	65
5	Evaluation . . . . .	67
5.1	Microbenchmark: Echo Throughput . . . . .	68
5.2	Microbenchmark: Echo Latency . . . . .	70
5.3	Application: Key-Value Store . . . . .	71
5.4	Application: Single Source Shortest Path (SSSP) . . . . .	75
6	Conclusion . . . . .	77
	Chapter 5 Conclusion and Future Work . . . . .	78
1	Conclusion . . . . .	78
2	Future Work . . . . .	78
2.1	Better Profiling Tools . . . . .	78
2.2	Better Load Balancing . . . . .	79
2.3	Handling Network Conditions . . . . .	79
	Bibliography . . . . .	81
	VITA . . . . .	89

## List of Tables

1	Minimal multi-threaded <i>Trust&lt;T&gt;</i> example. Reference counting ensures that the property remains in memory until the last <i>Trust&lt;T&gt;</i> referencing the property drops.	4
2	<i>Trust&lt;T&gt;</i> vs <i>Gossamer</i> code differences	6
3	Minimal <i>Trust&lt;T&gt;</i> example. An entrusted counter, referenced by <code>ct</code> is initialized to 17, then incremented once. The comments on the right indicate the types of the variables.	16
4	Asynchronous version of the example in Table 3. The second closure runs on the client, once the result of the first closure is received from the trustee.	19
5	Operations supported by each transport type.	54
6	Auto traits combined with negative trait implementations allow functions to restrict closures that capture pointers.	57
7	<i>Gossamer</i> disallows pointers from being captured and uses serialization to get around the limitation.	58

## List of Figures

1	Nested Delegation can lead to deadlocks. . . . .	21
2	Operation of <code>launch()</code> vs <code>apply()</code> . <code>launch()</code> supports blocking calls, including nested delegation calls in the delegated closure, but incurs a higher minimum overhead. Solid arrows indicate requests, dotted arrows are delegation responses. . . . .	22
3	The fixed-size $Trust<T>$ request slot consists of a <code>ready</code> bit, a request counter, and a variable number of variable-sized requests. The response slot contains a matching <code>ready</code> bit, as well as one (fixed, variable, or zero-sized) response per request in the matching request slot. There is one dedicated pair of request/response slots for each trustee/client pair. . . . .	27
4	Fetch-and-add throughput vs. object count. Uniform access distribution. . . . .	31
5	Fetch-and-add throughput vs. object count. Zipfian access distribution, $\alpha = 1$ . . . . .	31
6	Mean latency vs. offered load. Uniform access distribution. . . . .	34
7	Mean latency vs. offered load. Zipfian access distribution, $\alpha = 1$ . . . . .	34
8	Key-value store throughput, with 5% writes and varying table size. Uniform access distribution . . . . .	37
9	Key-value store throughput, with 5% writes and varying table size. Zipfian access distribution . . . . .	37
10	Key-value store throughput, with varying write percentage. Uniform access distribution . . . . .	38
11	Key-value store throughput, with varying write percentage. Zipfian access distribution . . . . .	38
12	Memcached throughput with varying table size. Uniform access distribution. S: stock memcached. . . . .	44
13	Memcached throughput with varying table size. Zipfian access distribution. . . . .	45
14	The throughput of key-value stores using compared systems with different key space size (5% write) . . . . .	47
15	To communicate using RDMA, the applications register some memory. The NICs can then read it via DMAs as necessary. . . . .	52
16	High-level design of <i>Gossamer</i> . . . . .	54
17	A <i>Gossamer</i> process on one machine . . . . .	55
18	Trustee is a worker and clients can delegate work using a trust that holds information about property and trustee . . . . .	57

19	NoRef trait forces a compile-time error message instead of run-time undefined behavior, if pointers are captured by a delegated closure. . . . .	58
20	<i>Gossamer</i> Request and Response slot layout . . . . .	60
21	Memory layout for Request Slots in client and trustee memory for a rack of three machines where each machine has three system threads. . . . .	64
22	Quadratic vs Reduced connections between threads. . . . .	66
23	Throughput of RDMA writes in Millions of operations per second (MOPs) for quadratic and reduced number of QPs. Each write is 256 bytes. . . . .	67
24	Throughput across eRPC, MPI, and <i>Gossamer</i> in an Echo application. Each request here is 32 bytes. . . . .	68
25	Throughput across a range of request sizes for eRPC, MPI, and <i>Gossamer</i> in an Echo application. This experiment uses 192 threads across 12 machines. . . . .	69
26	Throughput vs latency under load for <i>Gossamer</i> and eRPC in an Echo application with two machines and 16 threads per machine. . . . .	70
27	Throughput of key-value stores on up to 8 server clusters and 100% read workload with uniform distribution. . . . .	71
28	Throughput of key-value stores on up to 8 server clusters and 5% write workload with uniform distribution. . . . .	72
29	Throughput of key-value stores on up to 8 server clusters and 100% read workload with Zipfian distribution. . . . .	72
30	Throughput of key-value stores on up to 8 server clusters and 5% write workload with Zipfian distribution. . . . .	73
31	Time taken for SSSP on a graph with $2^{27}$ nodes and 32 average degree. . . . .	75
32	Time taken for SSSP on graphs with $2^n$ nodes and 32 average degree. . . . .	76

## Summary

As workloads have grown more parallel and data-heavy, single-threaded performance has become inadequate. This gives rise to the problem of managing safe access to shared data. Conventionally, this is achieved through locking or compare and swap, both of which use atomic instructions internally. As these instructions cannot be reordered with other instructions, they offer limited scaling, especially as the data becomes more contended. Delegation provides an alternative that scales well and is not as prone to high contention. Delegation is a versatile mechanism and can also be used to scale an application beyond a single machine. We design and implement two systems, *Trust*< $T$ > and *Gossamer*, to study how delegation compares against state-of-the-art systems in shared memory and distributed systems respectively.

## CHAPTER 1

### Introduction

Delegation is a concurrency technique where a thread hands off work on a shared object to a single designated owner instead of accessing the object directly. Because only the owner thread ever mutates the object, no locks are needed, other threads interact with it by sending requests. This avoids data races, simplifies reasoning, and can outperform lock-based designs by eliminating contention and complex synchronization. As delegation does not rely on shared memory, its applications extend naturally beyond local concurrency to broader distributed settings. We present two system, *Trust*<*T*> and *Gossamer*, to demonstrate the versatility of this approach across different domains.

*Trust*<*T*> introduces a type- and memory-safe alternative to traditional locking in concurrent programs. Instead of synchronizing multi-threaded access to an object of type *T* with a lock, it can be encapsulated in a *Trust*<*T*> with a designated trustee thread handling all operations through message passing. This delegation-based model eliminates per-lock throughput limitations, achieving up to 22x higher throughput in microbenchmarks and 5-9x improvement for a key-value store in high-contention workloads, while remaining competitive even without contention.

*Gossamer* extends the concept of delegation to distributed computing. It allows an application written for a single machine to scale to a small cluster with minimal changes. Using Remote Direct Memory Access (RDMA) for high-performance communication, *Gossamer* supports delegating Rust closures across machines and employs lightweight user-space

threads (fibers) for concurrency. Our experiments show 3x better throughput scaling than eRPC and performance comparable to the Graph500 MPI reference implementation.

Together, *Trust*< $T$ > and *Gossamer* illustrate how delegation unifies concurrency and distribution under a safe, efficient, and scalable programming paradigm, enabling high-performance computation without the complexities of shared memory or explicit synchronization.

The main contributions of this dissertation are summarized below:

- *Trust*< $T$ >: a model for efficient, multi-threaded, delegation-based programming with shared objects leveraging the Rust type system.
- *Gossamer*: an extension of *Trust*< $T$ > that enables a normal delegation based application to run on and utilize the resources of a rack with minimal changes to application code.

The rest of this chapter introduces both of these and then outlines the rest of this dissertation.

## 1 *Trust*< $T$ >

After single-threaded performance plateaued a decade or more ago, programmers have diverted their attention to making multi-threaded applications. This gives rise to the problem of synchronizing many threads as they access the same data for correctness and making it efficient to not lose too much performance when compared to single threaded applications. Today locks dominate this paradigm to the extent that there are hardware mechanisms in place to make the use of locking easier and more efficient. Hardware Lock Elision [6] and support for atomic increment instructions are examples of such mechanisms. However, a

delegation based programming model can offer better performance than a locking based one [57, 7, 42].

We present  $\text{Trust} < T >$  (pronounced *trust-tee*), a programming abstraction and runtime system which provides safe, high-performance access to a shared object (or **property**) of type  $T$ . Briefly, a  $\text{Trust} < T >$  provides a family of functions of the form:

$$\text{apply}(c : \text{FnOnce}(\&\text{mut } T) \rightarrow U) \rightarrow U,$$

which causes the closure  $c$  to be safely applied to the property (of type  $T$ ), and returns the return value (of type  $U$ ) of the closure to the caller. Here, **FnOnce** denotes a category of Rust closure types, and **&mut** denotes a mutable reference. (A matching set of non-blocking functions is also provided, which instead executes a callback closure with the return value.)

Critically, access to the property is only available through the  $\text{Trust} < T >$  API, which taken together with the Rust ownership model and borrow checker eliminates any potential for race conditions, given a correct implementation of *apply*. Our implementation of  $\text{Trust} < T >$  uses pure delegation. However, the design of the API also permits lock-based implementations, as well as hybrids.

Table 1a shows an example where a counter is incremented by two different threads. Here, the `clone()` call on `ct` (line 2) clones the trust, but not the property; instead a reference count is incremented for the shared property, analogous to `Arc::clone()`. On line 3, a newly spawned thread takes ownership of `ct2`, in the Rust sense of the word, then uses this to apply a closure (line 4). When the thread exits, `ct2` is dropped, decrementing the reference count, by means of a delegation request. When the last trust of a property

<pre> 1 let ct = local_trustee().entrust(17); 2 let ct2 = ct.clone(); 3 let thread = spawn(move    { 4   ct2.apply( c  *c+=1); 5 }); 6 ct.apply( c  *c+=1); 7 thread.join()?; 8 assert!(ct.apply( c  *c) == 19); </pre>	<pre> 1 let cm = Arc::new(Mutex::new(17)); 2 let cm2 = cm.clone(); 3 let thread = spawn(move    { 4   *(cm2.lock()?) += 1; 5 }); 6 *(cm.lock()?) += 1; 7 thread.join()?; 8 assert!(*cm.lock()? == 19); </pre>
---	---

(A) Example using *Trust*<*T*>.

(B) The same program using standard Rust primitives.

TABLE 1. Minimal multi-threaded *Trust*<*T*> example. Reference counting ensures that the property remains in memory until the last *Trust*<*T*> referencing the property drops.

is dropped, the property is dropped as well. For readers unfamiliar with Rust, Table 1b illustrates the rough equivalent of Table 1a, but using conventional Rust primitives instead of *Trust*<*T*>. Note the similarity in terms of legibility and verbosity.

Beyond the API, *Trust*<*T*> provides a runtime for scheduling request transmission and processing, as well as lightweight user threads (**fibers** below). This allows each OS thread to serve both as a Trustee, processing incoming requests, and a client. Multiple outstanding requests can be issued either by concurrent synchronous fibers or an asynchronous programming style. *Trust*<*T*> achieves performance improvements up to 22× vs. the best locks on congested micro-benchmarks and up to 9× on benchmarking workloads vs. stock `memcached`.

## 2 Gossamer

Multi-threaded applications perform much better than single threaded applications but are also more complex to design. The same is true when scaling applications from a single machine to a distributed setting. Here, lack of shared memory means that not all of the

compute units (threads/processes) can access all of the shared objects, which also complicates the synchronization mechanisms. Distributed applications, historically, rely on message passing to share data. Remote Procedural Calls (RPCs) is one such framework where threads can send requests to where data is located and receive responses upon completion [13, 62, 12, 60, 5, 65, 20, 14, 59, 45, 61]. This is very similar to how delegation works.

*Gossamer* extends  $Trust\langle T \rangle$  to build a rack-wide programming model. *Gossamer* aims to provide a high performance and easy to program in, framework that can take advantage of the increased resources available in a rack. Using the API provided by *Gossamer* programmers can code their applications like a normal delegation application with very few changes and the applications have the illusion of running on a single machine. Most of the distribution of work is handled behind the scenes, but programmers have the option to customize where the worker threads should run and where data should be held in the rack if they so choose. Table 2 shows a minimal  $Trust\langle T \rangle$  example and how the same program looks like when using *Gossamer*.

A key challenge that limits performance in distributed systems is the additional latency introduced by network communication. *Gossamer* mitigates this by combining RDMA with fibers. RDMA offers sub-microsecond one-way latency, significantly faster than accessing persistent storage on a single machine, though still slower than local memory and considerably slower than CPU cache accesses. Fibers are light weight user threads that run on top of system threads and help bridge this latency gap by allowing applications to issue and manage a large number of concurrent, outstanding requests. This is further amplified by batching the requests generated by all the fibers on a single system thread. While RDMA reduces latency, it does not scale efficiently as the number of connections per machine grows.

```

1 fn main() {
2     let pool = Threadpool::new(4);
3     let property = HashMap::new();
4     let trust = pool.entrust(property);
5     pool.spawn(|| {
6         trust.apply(|property| property.insert("Greeting", "Hello"));
7     });
8 }
```

(A) A small example of delegation code with *Trust<T>*

```

1 fn main() {
2     gossamer::main(gsm_main, 4);
3 }
4
5 fn gsm_main() {
6     let property = HashMap::new();
7     let trust = gossamer::entrust(property);
8     gossamer::spawn(|| {
9         trust.apply(|property| property.insert("Greeting", "Hello"));
10    });
11 }
```

(B) Same delegation request with *Gossamer*

TABLE 2. *Trust<T>* vs *Gossamer* code differences

To address this, *Gossamer* maintains only a single RDMA connection per hardware thread for each remote machine.

### 3 Dissertation Organization

The remainder of this dissertation is organized as follows. Chapter 2 gives brief background on shared memory synchronization and distributed systems. Chapter 3 introduces *Trust<T>*, a delegation based programming abstraction that provides safe access to shared data. Chapter 4 presents *Gossamer*, an extension of *Trust<T>*, that allows single-machine,

## Chapter 1. Introduction

multi-threaded applications built using *Trust*< $T$ >, to run on a rack consisting of many machines with minimal changes. Chapter 5 discusses future work and provides a conclusion.

## CHAPTER 2

### State of the Art

In this section, we provide background on the current landscape of shared-memory and distributed systems and explain how our proposed solutions advance the state of the art.

#### 1 Synchronizing shared memory systems

Locking suffers from a well-known scalability problem: as the number of contending cores grows, cores spend more and more of their time in contention, and less doing useful work. Consider a classical, but idealized lock, in which there are no efficiency losses due to contention. Here, the *sequential* cost of each critical section is the sum of (a) any wait for the lock to be released, (b) the cost of acquiring the lock, (c) executing the critical section, and (d) releasing the lock. Not counting any re-acquisitions on the same core, this must be at minimum one cache miss per critical section, in sequential cost. To make matters worse, this cache miss is incurred by an atomic instruction, effectively stalling the CPU until the cache miss is resolved (and in the case of a spinlock, until the lock is acquired).

Three main solutions to this problem exist. First, where the data structure permits, fine-grained locking can be used to split the data structure into multiple independently locked objects, thus increase parallelism, reduce lock contention and wait times. With the data structure split into sufficiently many objects, and *accesses distributed uniformly*, a fine-grained locking approach tends to offer the best available performance.

The second solution is to use of nonblocking algorithms [57, 28, 48, 8, 47, 9, 37, 38].

Nonblocking algorithms handle mutual exclusion by ensuring progress through atomic operations, allowing threads to proceed without waiting for others to release a lock. Compared to locks, they offer superior scalability under contention, eliminate deadlock and priority inversion, and often provide better worst-case latency, which makes them attractive for highly concurrent or real-time systems. However, they are not as widely used as locks because they are significantly harder to design, reason about, and debug, with subtle correctness issues such as ABA problems and complex memory-reclamation requirements.

The third solution is various forms of delegation, where one thread has custody of the object, and applies critical sections on behalf of other threads. Ideally, this minimizes the sequential cost of each critical section without changing the data structure: there are no sequential cache misses, ideally no atomic instructions, but of course the critical sections themselves still execute sequentially.

*Combining* [27, 32, 22, 26, 50, 71, 58], is a flavor of delegation in which threads temporarily take on the role of *combiner*, performing queued up critical sections for other threads. Combining can scale better than locking in congested settings, but does not offer the full benefits of delegation as it makes heavy use of atomic operations, and moves data between cores as new threads take on the *combiner* role. Most recently, TClocks [31] offers a fully transparent combining-based replacement for locks, by capturing and restoring register contents, and automatically pre-fetching parts of the stack. TClocks claims substantial benefits for extremely congested locks, and the backward compatibility is of course quite attractive. However, a cursory evaluation in Section 6 reveals that TClocks substantially

underperform regular locks beyond extremely high contention settings, and never approaches  $Trust < T >$  performance.

Beyond *combining*, delegation has primarily been explored in proof-of-concept or one-off form, with relatively immature programming abstractions. We propose  $Trust < T >$ , a full-fledged delegation API for the Rust language, which presents delegation in a type-safe and familiar form, while substantially outperforming the fastest prior work on delegation.

While delegation offers much higher throughput for congested shared objects, it does suffer higher latency than locking in uncongested conditions. To hide this latency, and make delegation competitive in uncongested settings,  $Trust < T >$  exposes additional concurrency to the application via asynchronous delegation requests and/or light-weight, delegation-aware user threads (*fibers*).

Lacking modularity is another common criticism of delegation: in FFWD [57], an early delegation design, delegated functions must not perform any blocking operations, which includes any further delegation calls. In  $Trust < T >$ , this constraint remains for the common case, as this typically offers the highest efficiency. However,  $Trust < T >$  offers several options for more modular operation. First, asynchronous/non-blocking delegation requests are not subject to this constraint - these requests may be safely issued in any context. Second, leveraging our light-weight user threads, we offer the option of supporting blocking calls in delegated functions, on an as-needed basis.

Finally, prior work on delegation has required one or more cores to be dedicated as delegation servers. While  $Trust < T >$  offers dedicated cores as one option, the  $Trust < T >$  runtime has every core act as a delegation server, again leveraging light-weight user threads. Beyond easing application development and improving load balancing, having a delegation server on

every core allows us to implement  $Trust\langle T \rangle$  without any use of atomic instructions, instead relying on delegation for all inter-thread communication. Beyond potential performance advantages, this also makes  $Trust\langle T \rangle$  applicable to environments where atomic operations are unavailable.

## 2 Distributed Systems

Just as multithreaded applications exploit multiple cores within a single machine, distributed systems exploit resources across multiple machines. To support this model, a variety of distributed systems frameworks have emerged, each providing abstractions and tooling to handle communication, coordination, and fault tolerance. Below we discuss a few of these briefly.

Remote Procedure Call (RPC) is a widely used model for communication in datacenters [13, 62, 12, 60, 5, 65, 20, 14, 59, 45, 61]. RPCs provide the useful abstraction of running a function on a remote machine, while hiding the underlying network communication. For most approaches to RPC, the function needs to be registered at the start and calls to it can be either synchronous or asynchronous. A synchronous call will wait for a response while an asynchronous call can provide a callback function to be executed once the response is ready and continue execution. This is similar to delegation, rather than moving data to where it is needed, the code that uses the data is sent to its location. This process typically involves serializing and deserializing of data and transport protocols that ensure reliable message delivery. By decoupling distributed computation from explicit message handling, RPC simplifies the development of networked applications, supports modular system design,

and serves as a foundational mechanism in many modern distributed architectures, including microservices and cloud-based systems.

Message Passing is another dominant communication framework [17, 24, 16, 15, 52]. Unlike RPC, Message Passing has explicit communication and involves sending and receiving messages. Depending on the implementation, there can be guarantees regarding the reliability of message transport or the order of messages, or it can be a send-and-forget system where the sender might not even know if the message reached its destination. MPI [29] a widely used standard that provides a range of communication primitives, from one-to-one messages to many-to-many. It also provides both blocking and non blocking versions of send/receive.

There is also ongoing research to minimize resource usage of communication as much as possible. Remote Direct Memory Access (RDMA) is one such attempt [54, 56, 64, 53]. RDMA allows the user to read a remote machine's memory by utilizing specialized Network Interface Cards (NIC) while bypassing the remote machine's operating system. Additionally, there is also some work on using a custom networking stack for distributed systems as compared to a hardware accelerated RDMA system [35, 44, 10].

RDMA has also given rise to a renewed academic interest in distributed shared memory (DSM) [67, 72, 30, 41, 68, 43]. The idea behind DSM is to create the illusion of a shared address space for networked applications. The application can be programmed as if all of the data is located in local memory and the DSM system is responsible for ensuring safety, consistency and synchronization. One challenge with DSM is the coherence overhead, which can dominate both CPU cycles and network traffic, increasing latency [43]. [68] also

highlights the need to revisit cache and buffer management due to narrowing gap between local and remote memory.

*Gossamer* combines RPC and some aspects of DSM to simplify the synchronization of remote memory while delivering high performance. It extends the delegation mechanism from *Trust*< $T$ > to work with distributed memory enabling distributed applications to be designed like shared memory applications with a few limitations. As only a single thread ever accesses any piece of shared data, *Gossamer* doesn't need to do any extra work to ensure safety and consistency of memory. Along with simplifying the design, this also reduces data access overhead lowering the latency.

## CHAPTER 3

### ***Trust*<*T*>**

#### 1 Summary

*Trust*<*T*> is a general, type- and memory-safe alternative to locking in concurrent programs. Instead of synchronizing multi-threaded access to an object of type *T* with a lock, the programmer may place the object in a *Trust*<*T*>. The object is then no longer directly accessible. Instead, a designated thread, the object’s *trustee*, is responsible for applying any requested operations to the object, as requested via the *Trust*<*T*> API.

While locking offers a limited throughput *per lock*, *Trust*<*T*> is based on delegation, a message-passing technique which does not suffer this per-lock limitation. Instead, per-object throughput is limited by the capacity of the object’s trustee, which is typically considerably higher.

Our evaluation shows *Trust*<*T*> consistently and considerably outperforming locking where lock contention exists, with up to 22× higher throughput in microbenchmarks, and 5–9 × for a home grown key-value store, in situations with high lock contention. Moreover, *Trust*<*T*> is competitive with locks even in the absence of lock contention.

#### 2 Introduction

Safe access to shared objects is fundamental to many multi-threaded programs. Conventionally, this is achieved through *locking*, or in some cases through carefully designed lock-free data structures, both of which are implemented using atomic compare-and-swap

(CAS) operations. By their nature, atomic instructions do not *scale* well: atomic instructions must not be reordered with other instructions, often starving part of today’s highly parallel CPU pipelines of work until the instruction has retired. This effect is exacerbated when multiple cores are accessing the same object, resulting in the combined effect of frequent cache misses and cores waiting for each other to release the cache line in question, while the atomic instructions prevent them from doing other work. These effects are further strengthened in NUMA architectures [36, 70, 66, 51].

Delegation [22, 18, 55, 26, 32, 50, 71, 58, 21, 57], also known as light-weight remote procedure calls (LRPC), offers a highly scalable alternative to locking. Here, each shared object<sup>1</sup> is placed in the care of a single core. Other cores issue requests to this designated core, specifying operations to be performed on the object.

Compared to locking, where threads typically contend for access, and may even suspend execution to wait for access, delegation requests from different clients are submitted to the trustee in parallel and without contention. This dramatically reduces the cost of coordination for congested objects. The operations/critical sections are applied sequentially in both designs: by each thread using locks, or by the trustee using delegation; here delegation may benefit from improved locality at the trustee. Together, this translates to much higher maximum per-object throughput with delegation vs. locking.

However, under medium or low contention, classical delegation struggles to compete with locking: the latency and overhead of request transmission, request processing and response transmission are insignificant compared to the cost of contending for a lock, but can be substantial compared to the cost of acquiring an *uncontended* lock.

---

<sup>1</sup>Here, we use *object* to mean a data structure that would be protected by a single lock.

### 3 *Trust*<*T*>: The Basics

```

1 let ct = local_trustee().entrust( 17 );           // ct: Trust<i32>
2 ct.apply( |c| *c+=1 );                           // c: &mut i32
3 assert!(ct.apply( |c| *c ) == 18 );

```

TABLE 3. Minimal *Trust*<*T*> example. An entrusted counter, referenced by *ct* is initialized to 17, then incremented once. The comments on the right indicate the types of the variables.

The objective of *Trust*<*T*> is to provide an intuitive API for safe, efficient access to shared objects. Naturally, our design motivation is to support delegation, but the *Trust*<*T*> API can in principle also be implemented using locking, or a combination of locking and delegation. Below, we first introduce the basic *Trust*<*T*> programming model, as well as the key terms *trust*, *property*, *trustee* and *fiber* in the *Trust*<*T*> context, before digging deeper into the design of *Trust*<*T*>.

#### 3.1 Trust: a reference to an object

A *Trust*<*T*> is a thread-safe reference counting smart-pointer, similar to Rust’s *Arc*<*T*>. To create a *Trust*<*T*>, we clone an existing *Trust*<*T*> or *entrust* a new object, or *property* of type *T*, that is meant to be shared between threads. Once entrusted, the property can only be accessed by *applying* closures to it, using a trust. Table 3 illustrates this through a minimal Rust example. Line 1 entrusts an integer, initialized to 17, to the local trustee - the trustee fiber running on the current kernel thread. Line 2 applies an anonymous closure to the counter, via the trust. The closure expected by *apply* takes a mutable reference to the property as argument, allowing it unrestricted access to the property, in this case, our integer. The example closure increments the value of the integer. The assertion on line 3

is illustrative only. Here, we apply a second closure to retrieve the value of the entrusted integer<sup>2</sup>.

### 3.2 Trustee - a thread in charge of entrusted properties

In our examples above, *Trust<T>* is implemented using delegation. Here, a *property* is *entrusted* to a *trustee*, a designated thread which executes applied closures on behalf of other threads. In the default *Trust<T>* runtime environment, every OS thread in use already has a trustee user-thread (*fiber*) that shares the thread with other fibers. When a fiber applies a closure to a trust, this is sent to the corresponding trustee as a message. Upon receipt, the trustee executes the closure on the property, and responds, including any closure return value. This may sound complex, yet the produced executable code substantially outperforms locking in congested settings.

A `TrusteeReference` API is also provided. Here, the most important function is `en-trust()`, which takes a property of type `T` as argument (by value), and returns a *Trust<T>* referencing the property that is now owned by the trustee. This API allows the programmer to manually manage the allocation of properties to trustees, for performance tuning or other purposes. Alternatively, a basic thread pool is provided to manage distribution of fibers and variables across trustees.

### 3.3 Fiber - a delegation-aware, light-weight user thread

While the *Trust<T>* abstraction has some utility in isolation, it is most valuable when combined with an efficient message-passing implementation and a user-threading runtime.

---

<sup>2</sup>A note on ownership: While the passed-in closure takes only a reference to the property, the Rust syntax `*c` denotes an explicit dereference, essentially returning a copy of the property to the caller. This will pass compile-time type-checking only for types that implement `Copy`, such as integers.

User-level threads, also known as coroutines or *fibers*, share a kernel thread, but each execute on their own stack, enabling a thread to do useful work for one fiber while another waits for a response from a trustee. This includes executing the local trustee fiber to service any incoming requests.

In this default setting, the synchronous `apply()` function suspends the current fiber when it issues a request, scheduling the next fiber from the local ready queue to run instead. The local fiber scheduler will periodically poll for responses to outstanding requests, and resume suspended fibers as their blocking requests complete.

### 3.4 Delegated context

For the purpose of future discussion, we define the term *delegated context* to mean the context where a delegated closure executes. Generally speaking, closures execute as part of a trustee fiber, on the trustee's stack. Importantly, blocking delegation calls are not permitted from within delegated context, and will result in a runtime assertion failure. In Section 4, we describe multiple ways around this constraint.

## 4 Core API

The *Trust< T >* API supports a variety of ways to delegate work, some of which we elide due to space constraints. Below, we describe the core functions in detail.

### 4.1 `apply(): synchronous delegation`

```
apply(c: FnOnce(&mut T)->U)->U
```

`apply()` is the primary function for blocking, synchronous delegation as described in earlier sections. It takes a closure of the form `|&mut T| {}`, where `T` is the type of the property. If the closure has a return value, `apply` returns this value to the caller.

Importantly, `apply()` is synchronous, suspending the current fiber until the operation has completed. Often, the best performance with `apply()` is achieved when running multiple application fibers per thread. Then, while one fiber is waiting for its response, another may productively use the CPU.

## 4.2 apply\_then(): non-blocking delegation

```
apply_then(c: FnOnce(&mut T)->U,
           then: FnOnce(U))
```

```
1 let ct = trustee.entrust(17);           // create trust for shared counter set to 17
2 ct.apply_then(|c| { *c+=1; *c },        // increment counter and return its value
3                 |val| assert!(val==18)); // check return value once received
```

TABLE 4. Asynchronous version of the example in Table 3. The second closure runs on the client, once the result of the first closure is received from the trustee.

Frequently, asynchronous (or non-blocking) application logic can allow the programmer to express additional concurrency either without running multiple fibers, or in combination with multiple fibers. Here, `apply_then()` returns to the caller without blocking, and does not produce a return value. Instead, the second closure, `then`, is called with the return value from the delegated closure, once it has been received. Table 4 demonstrates the use of `apply_then()` following the pattern of Table 3.

The `then`-closure is a very powerful abstraction, as it too is able to capture variables from the local environment, allowing it to perform tasks like adding the return value (once available) to a vector accessible to the caller. Here, Rust's strict lifetime rules automatically catch otherwise easily introduced use-after-free and dangling pointer problems, forcing the programmer to appropriately manage object lifetime either through scoping or reference counted heap storage.

Importantly, as `apply_then()` does not suspend the caller, it may freely be called from within delegated context.

#### 4.3 `launch()`: apply in a trustee-side fiber

```
launch(c: FnOnce(&mut T)->U)->U
launch_then(c: FnOnce(&mut T)->U,
            then: FnOnce(U))
```

The most significant constraint imposed by *Trust*<*T*> on the closure passed to `apply()` and `apply_then()` is that the closure itself may not block. Blocking in delegated context means putting the trustee itself to sleep, preventing it from serving other requests, potentially resulting in deadlock as shown in Figure 1. In previous work [42], this problem was addressed by maintaining multiple server OS threads, and automatically switching to the next server when one server thread blocks. This avoids blocking the trustee, but imposes high overhead, resulting in considerably lower performance, as demonstrated in [57].

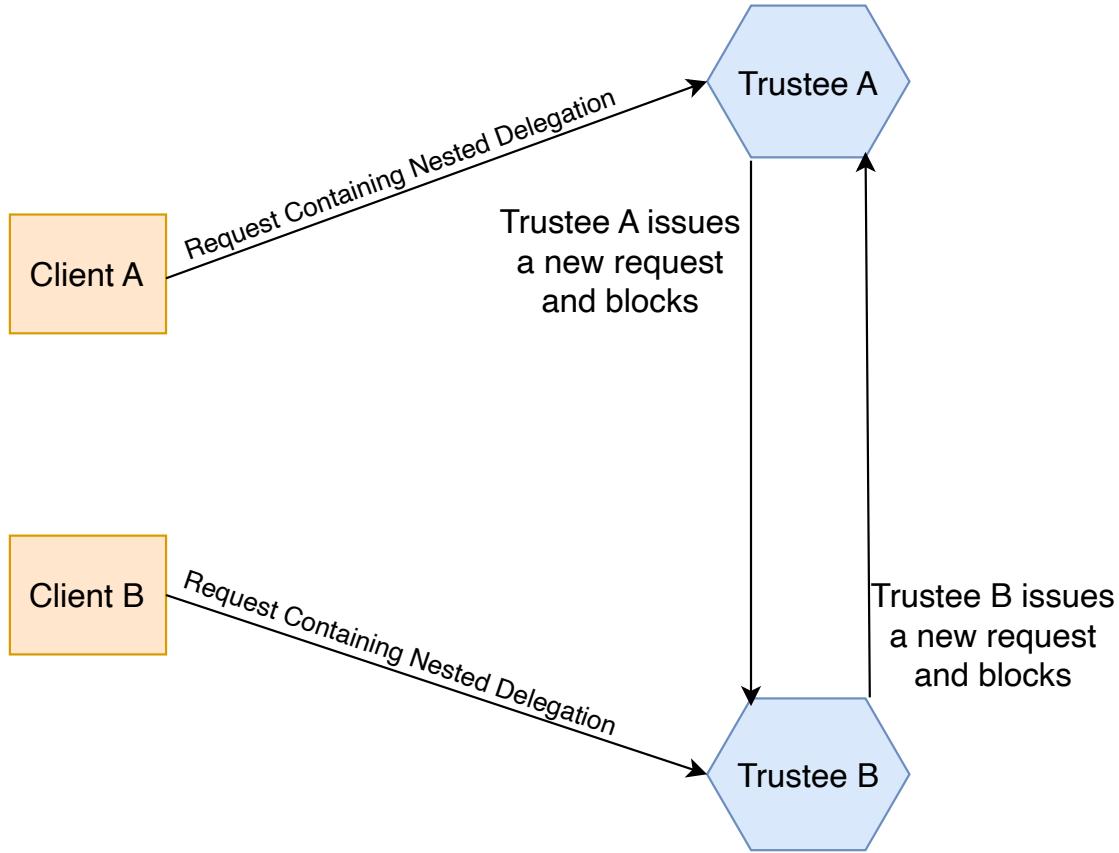


FIGURE 1. Nested Delegation can lead to deadlocks.

In *Trust*<*T*>, blocking in delegated context is prohibited: attempted suspensions in delegated context are detected at runtime, resulting in an assertion failure. Closures may still use `apply_then()`, but not the blocking `apply()`.<sup>3</sup>

The lack of *nested blocking delegation* can be a significant constraint on the developer, and perhaps the most important limitation of *Trust*<*T*>. Specifically, it affects modularity, as a library function that blocks internally, even on delegation calls, cannot be used from within delegated context.

<sup>3</sup>Other forms of blocking, such as I/O waits or scheduler preemption, do not result in assertion failures. However, these can significantly impact performance if common, as blocking the trustee can prevent other threads from making progress.

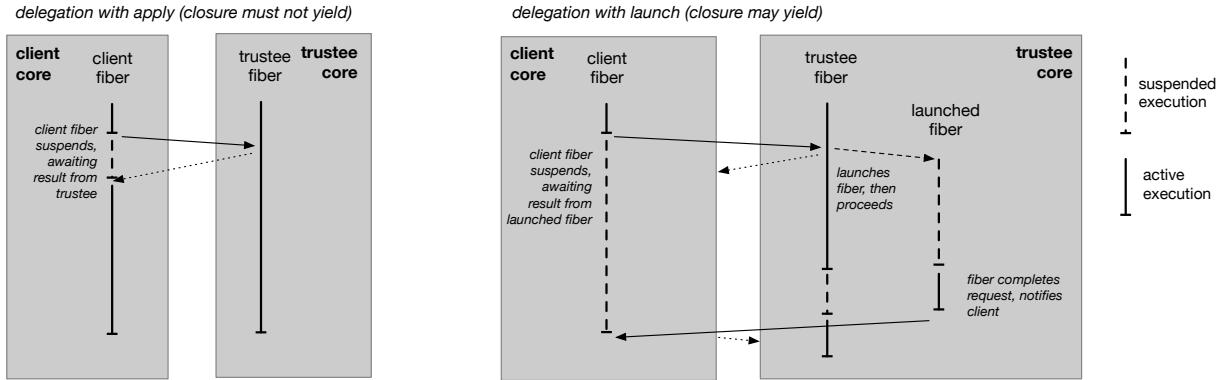


FIGURE 2. Operation of `launch()` vs `apply()`. `launch()` supports blocking calls, including nested delegation calls in the delegated closure, but incurs a higher minimum overhead. Solid arrows indicate requests, dotted arrows are delegation responses.

To address this, without sacrificing the performance of the more common case, we provide a convenience function: `launch()`, which offers all the same functionality as `apply()`, but without the blocking restriction. Figure 2 describes `launch()` from an implementation standpoint. `launch()` creates a temporary fiber on the trustee's thread, which runs the closure. If this fiber is suspended, the client is notified, and the trustee continues to serve the next request. Once the temporary fiber resumes and completes execution of the closure, it then delivers the return value and resumes the client fiber via a second delegation call. Thus, if a delegated closure fails the runtime check for blocking calls, the developer can fix this by replacing the `apply()` call, with a `launch()` call.

#### 4.3.1 Atomicity and `launch()`

That said, a complicating factor with blocking closures executed by `launch()` is that without further protection, property accesses are no longer guaranteed to be atomic: while the newly created fiber is suspended, another delegation request may be applied to the property, resulting in a race condition. To avoid this risk, `launch()` is implemented only for

`Trust<Latch<T>>`. `Latch<T>` is a wrapper type which provides mutual exclusion, analogous to `Mutex<T>` except that it uses no atomic instructions, and thus may only be accessed by the fibers of a single thread.<sup>4</sup>

#### 4.3.2 Leveraging Rust for safe and efficient delegation

Using the Rust type system, we ensure that delegated closures in `Trust<T>` cannot capture values that contain any references or pointers. This is discussed in more detail in Section Figure 4.4.

In principle, this is far stricter than what is necessary: the existing and pervasive Rust traits `Send` and `Sync` already describe the types that may be safely moved and shared between threads, and this continues to hold within `Trust<T>`.

That said, safety does not guarantee performance. A common performance pitfall when writing delegation-based software is memory stalls on the trustee, which affects trustees disproportionately due to the polling nature of the delegation channel (see Section 5.1). Frequent cache misses and use of atomic instructions in delegated closures can substantially degrade trustee throughput vs. running closures with good memory locality.

Generally speaking, cache line contention and use of atomic instructions are a natural result of sharing memory between threads. By prohibiting the capture of references and pointers, `Trust<T>` makes accidental shared memory patterns of programming much less likely in delegated code, and encourage cache friendly pass-by-value practices.

---

<sup>4</sup>In Rust terms, `Latch<T>` does not implement `Sync`.

### 4.3.3 Variable-size and other heap-allocated values

Rust closures very efficiently and conveniently capture their environment, which `apply()` sends whole-sale to the trustee. However, only types with a size known at compile time may be captured in a Rust closure (or even allocated on the stack).

In conventional Rust code, variable size types, including strings, are stored on the heap, and referenced by a `Box<T>` smart pointer. For the reasons described above (see Section 4.3.2), we do not allow `Box<T>` or other types that include pointers or references to be captured in a closure: only pure values may pass through the delegation channel.

As a result, variable size objects and other heap-allocated objects must be passed as explicit arguments rather than captured, so that they may be serialized before transmission over the delegation channel. For example, a `Box<[u8]>` (a reference to a heap-allocated variable-sized array of bytes) cannot traverse the delegation channel. Instead, we encode a copy of the variable number of bytes in question into the channel, and pass this value to the closure when it is executed by the trustee. In practice, this takes the form of a slightly different function signature.

```
apply_with(c: FnOnce(&mut T, V)->U, w: V)->U
```

Here, the `w:` argument is any type `V:Serialize+Deserialize`, using the popular traits from the `serde` crate. That is, any type that can be serialized and deserialized, may pass over the delegation channel in serialized form. If more than one argument is needed, these may be passed as a tuple. Thus, to insert a variable-size key and value into an entrusted table, we might use:

```
table_trust.apply_with(|table, (key, value)|  
    table.insert(key, value), (key, value))
```

We use the efficient `bincode` crate internally for serialization. As a result, while passing heap-allocated values does incur some additional syntax, the impact in terms of performance is minimal.

## 5 Key Design and Implementation Details

In this section, we delve deeper into the design and implementation of *Trust< T >*, from the mechanics of delegating Rust closures and handling requests and responses, to asynchronous versions of `apply()`.

### 5.1 Delegating Closures

The key operation supported by *Trust< T >* is `apply()`, which applies a Rust *closure* to the property referenced by the trust. A Rust closure consists of an anonymous function and a captured environment, which together is represented as a 128-bit *fat pointer*. Thus, to delegate a closure, a request must at minimum contain this fat pointer, and a reference to the property in question.

One or more requests are written to the client's dedicated, fixed-sized *request slot* for the appropriate trustee. That is, only the client thread may write to the request slot. For efficiency, if the captured environment of the closure fits in the request slot, we copy the environment directly to the slot, and update the fat pointer to reflect this change. A flag in the request slot indicates that new requests are ready to be processed. See Section 5.3 for details on request and response slot structure.

Responses are transmitted in a matching dedicated response slot. Leveraging the Rust type system, we restrict both requests and responses to types that can be serialized. The subtle implication of this is that the return value may not pass any references or pointers to trustee-managed data.<sup>5</sup> While small closures with simple, known-and-fixed-size return types will generally yield the best performance, there is no limit beyond the serializability requirement on the size or complexity of closures and return types.

## 5.2 Scheduling Delegation Work

Generally speaking, a call to `apply()` appends a request to a pending request queue, local to the requesting thread. In the case of `apply()`, the calling fiber is then suspended, to be woken up when the response is ready. Pending requests are sent during response polling, and as soon as an appropriate request slot is available. The intervening time is spent running other fibers, including the local trustee fiber, and polling for responses/transmitting requests.

There is a throughput/latency trade-off between running application fibers, and polling for requests/responses: poll too often, and few requests/responses will be ready, wasting polling effort. Poll too seldom, and many requests/responses will have been ready for a long time, increasing latency. Automatically tuning this trade-off is an area of ongoing research. That said, the current implementation performs delegation tasks in a fiber that is scheduled in FIFO order just as other fibers. After serving incoming requests, this fiber polls for incoming responses and issues any enqueued outgoing requests as applicable.

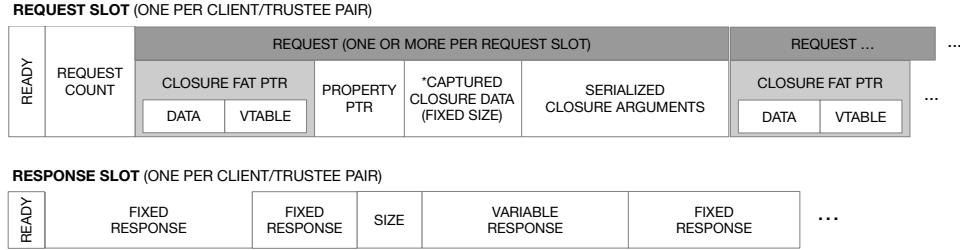


FIGURE 3. The fixed-size *Trust<T>* request slot consists of a `ready` bit, a `request count`, and a variable number of variable-sized requests. The response slot contains a matching `ready` bit, as well as one (fixed, variable, or zero-sized) response per request in the matching request slot. There is one dedicated pair of request/response slots for each trustee/client pair.

### 5.3 Request and Response Slot Structure

Figure 3 illustrates the internal structure of the basic request and response slot design.

A header consisting of a `ready` bit and a `request count`, is followed by a variable number of variable-sized requests. The value of the `ready` bit is used to indicate whether a new request or set of requests has been written to the slot: if the bit differs from the `ready` bit in the corresponding response slot, then a new set of requests is ready to be processed.

By default, the slot size is 1152 bytes, and the client may submit as many closures as it can fit within the slot. Here, the minimum size of a request is 24 bytes: a 128-bit fat pointer for the closure, and a regular 64-bit pointer for the property. The captured environment of Rust closures have a known, fixed size, which is found in the vtable of the closure. For typical small captured environments, this is copied into the request slot, and the pointer updated to point at the new location. Serialized closure arguments are appended next, followed by the next request.

---

<sup>5</sup>That said, we cannot prevent determined Rust programmers from using `unsafe` code to circumvent this restriction.

Responses are handled in a similar fashion, though there is no minimum response size.

Responses are sent simultaneously for all the requests in the request slot. The size of each response is often statically known, in which case it is not encoded in the channel. Any variable-size responses are preceded by their size.

The size of each request is always known, either statically or at the time of submission, which means we can restrict the number of requests sent to what can be accommodated by the request slot. The size of the response is not always known at the time the request is sent. In cases where the combined size of return values exceeds the space in the response slot, the trustee dynamically allocates additional memory to fit the full set of responses, at a small performance penalty.

### 5.3.1 Two-part slot optimization

In order to accommodate a broad range of application characteristics, including those with a single trustee and many clients, as well as a single client with many trustees, we introduce a small optimization beyond the basic design above. Rather than represent the request and response slots as monolithic blocks of bytes, we represent each as two blocks: a 128-byte primary block, and a 1024-byte overflow block; each request and response is written, in its entirety, to one or the other block.

This addresses an otherwise problematic trade-off with respect to the request and response slot sizes: with a monolithic request slot of, say, one kilobyte, the trustee would be periodically scanning flags 1024 bytes apart, a very poor choice from a cache utilization perspective, unless the slots are heavily utilized. A two-part design accommodates a large

number of requests (where needed), but improves the efficiency of less heavily utilized request slots by spacing ready flags, and a small number of compact requests, more closely using a smaller primary request block.

#### 5.4 Reference counting for Trusts

Trusts in *Trust<T>* act as rust smart pointers that own the property, since they can be used to access and modify the property behind them. This means that we need to make sure that when all of the trusts are dropped, the property is also dropped and the associated memory is freed so as not to have memory leakage. To achieve that trusts need to be reference counted, but a naïve integer count will not suffice due to a combination of the following two issues.

- *Trust<T>* does not support a blocking delegation operation when in the middle of another delegation request. For example calling *apply* on a trust from within a closure that itself is used for delegation would cause the system to hang and never finish. In the implementation, a runtime assertion raises a panic if this condition is encountered.
- If the integer used for counting references, *refcount* below, is incremented and decremented asynchronously, i.e. with a nonblocking delegation call, it could lead to use after free bugs. Consider the following scenario: Thread A clones a trust with only one reference and sends an async request to increment its refcount. The cloned trust is then sent to Thread B that drops it, sending another async request to decrement its refcount. While the requests issued by the same fiber are guaranteed to be completed in the same order they are issued, there is no such guaranty across multiple

threads or even fibers. It is entirely possible that the decrement request is processed first, making the refcount zero, which results in the property being dropped. Thread A however still has a valid trust to this property which it can use, expecting the property to still be available.

Not being able to use async delegation to increment and decrement the refcount leads to not being able to clone a trust from within a delegated context, which can be quite restrictive. To get around that, *Trust<T>* uses a new way to keep track of how many trusts are active at any time. Each trust is given a unique id at the time of creation, be it a new trust or a clone of an existing one. Instead of just using an integer, *Trust<T>* uses a set of these ids associated with each property. Instead of incrementing or decrementing the refcount, clone and drop both issue a delegation request involving a symmetric difference operation. Symmetric difference is defined as adding an element to a set if it is not already a member, and removing it from the set if it is. This way, regardless of the order in which requests originating in clone and drop are processed, the first one will add the trust id to the set and later one will remove it, allowing us to use async delegation for both. This, in turn, enables us to clone a trust within a delegated context.

## 6 Evaluation

Below, we evaluate the performance of *Trust<T>* in two ways: 1) on both microbenchmarks, designed to stress test the core mechanisms behind *Trust<T>* and locking, and 2) on end-to-end application benchmarks, which measure the performance impact of *Trust<T>* in the context of a complete system and a more realistic use case.

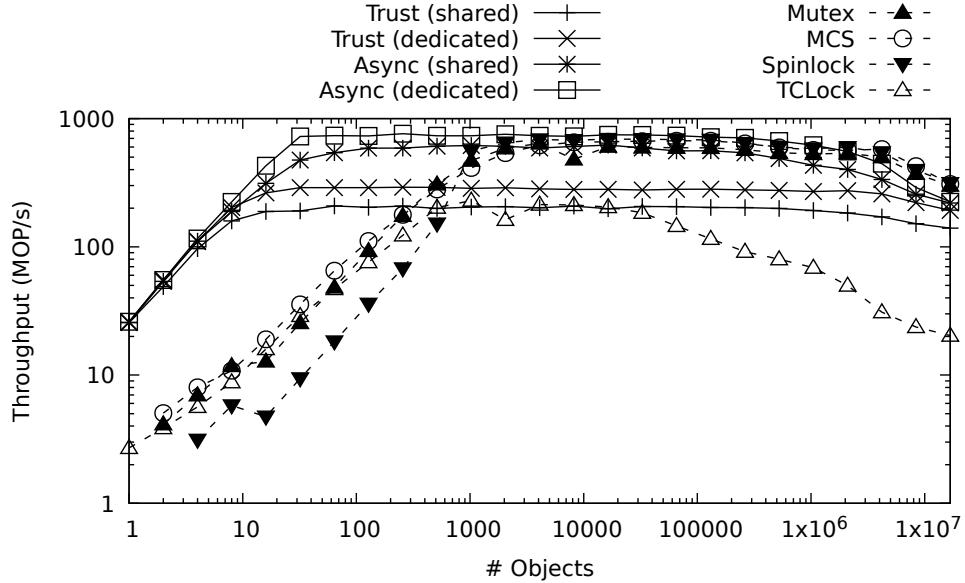


FIGURE 4. Fetch-and-add throughput vs. object count. Uniform access distribution.

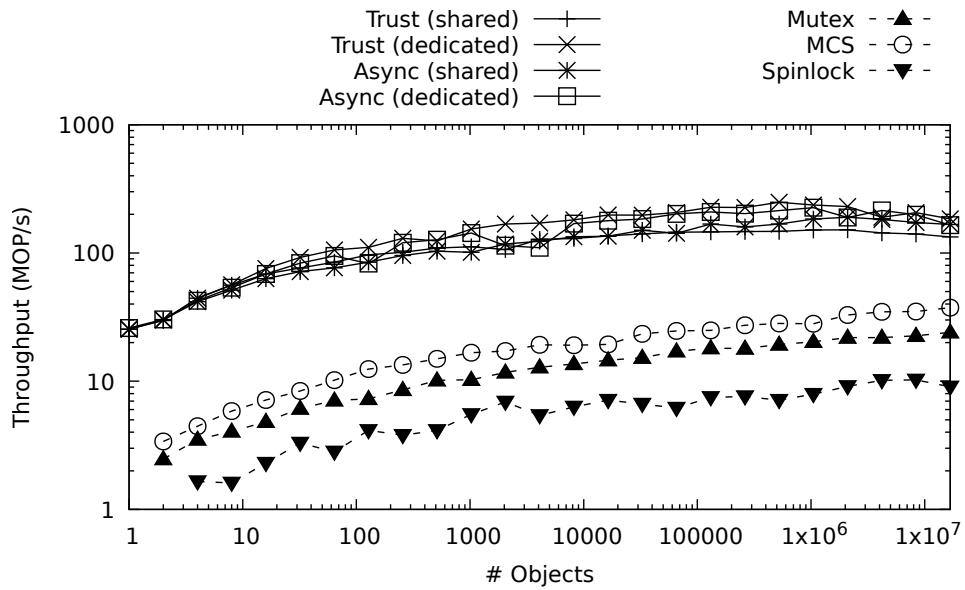


FIGURE 5. Fetch-and-add throughput vs. object count. Zipfian access distribution,  $\alpha = 1$ .

## 6.1 Fetch and Add: Throughput

For our first microbenchmark, we use a basic fetch-and-add application. Here, a number of threads repeatedly increment a counter chosen from a set of one or more, and fetches

the value of the counter. In common with prior work on synchronization and delegation [22, 18, 55, 26, 32, 50, 71, 58, 21], we also include a single pause instruction in both the critical section and the delegated closures. The counter is chosen at random, either from a uniform distribution, or a Zipfian distribution. Each thread completes 1 million such increments. In this section, each data point is the result of a single run.

Below, we primarily evaluate on a two-socket Intel Xeon CPU Max 9462, of the Sapphire Rapids architecture. This machine has a total of 64 cores, 128 hyperthreads, and 384 GB of RAM. Unless otherwise noted, we use 128 OS threads. In testing, several older x86-64 ISA processors have shown similar trends – these results are not shown here. For locking solutions, we use standard Rust `Mutex<T>` and the spinlock variant provided by the Rust `spin-rs-0.9.8` crate, as well as `MCSLock<T>` provided by the Rust `synctools-0.3.2` crate. For `Trust<T>`, we show results for blocking delegation (`Trust`) as well as nonblocking delegation (`Async`). In Figure 4, we also include TCLocks, a recent combining approach offering a transparent replacement for standard locks, via the Litl lock wrapper [33] for `pthread_mutex`. To be able to evaluate this lock, we wrote a separate C microbenchmark, matching the Rust version. In the interest of an apples-to-apples comparison, we first verified that the reported performance with stock `pthread_mutex` on the C microbenchmark matched the Rust `Mutex<T>` performance in our Rust microbenchmark.

Below, the `Trust` results may be seen to represent any application with ample concurrency available in the form of conventional synchronous threads. `Async` represents applications where a single thread may issue multiple simultaneously outstanding requests, e.g. a key-value store or web application server. Applications with limited concurrency are not well suited to delegation, except where the delegated work is itself substantial, which is not the

case for this fetch-and-add benchmark. We further report results with both letting all cores serve as both clients and trustees (**shared**), as well as with an ideal number of cores dedicated serve only as trustees (**dedicated**).

### 6.1.1 Uniform Access Pattern

Figure 4 illustrates the performance of several solutions on the uniform distribution version of this benchmark. For a very small number of objects, no data points are reported for some of the lock types - this is because the experiment took far too long to run due to severe congestion collapse.

`Trust<T>` substantially outperforms locking under congested conditions. Between 1–16 objects, the performance advantage is 8–22 $\times$  the best-performing MCSLock. For larger numbers of objects, the overhead of switching between fibers becomes apparent, as asynchronous delegation is able to reach a higher peak performance. In entirely uncongested settings, with 10 $\times$  as many objects as there are threads, locking is able to match asynchronous delegation performance. TCLocks [31] was the only lock type to complete the single-lock experiment within a reasonable time. It consistently outperforms spinlocks under congestion, and remains competitive with Mutex and MCS on highly congested locks. However, TCLocks appear to trade their transparency for high memory and communication overhead, making it unable to compete performance-wise beyond highly congested settings.<sup>6</sup> Moreover, we struggled to apply TCLocks to memcached (which consistently crashed under high load), as well as to Rust programs (as Rust now uses built-in locks rather than `libpthread` wrappers). We thus elide TCLocks from the remainder of the evaluation.

---

<sup>6</sup>TCLocks performance appears somewhat architecture dependent. In separate runs on our smaller Skylake machines, TCLocks were able to outperform Mutex by  $\approx$ 50% under the most extreme contention (a single lock).

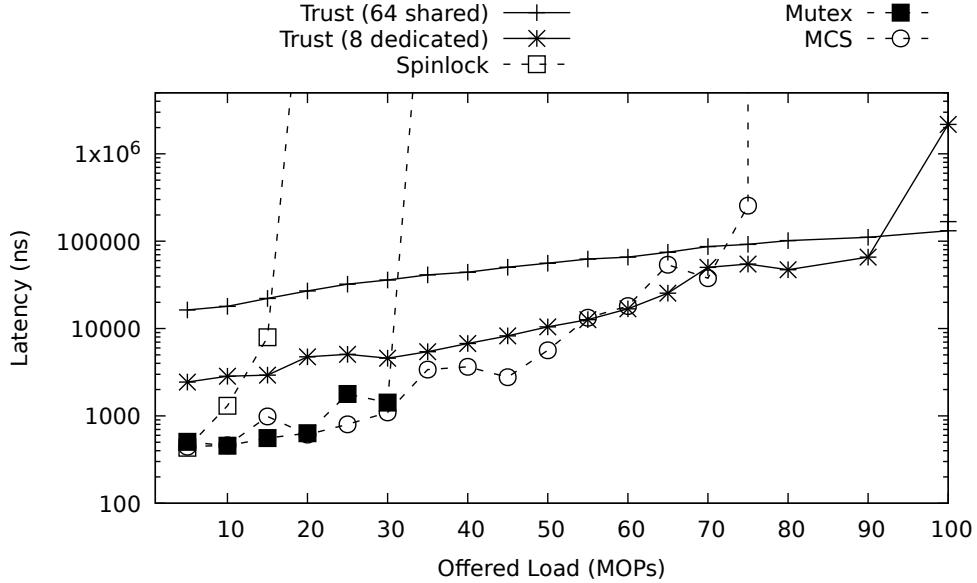
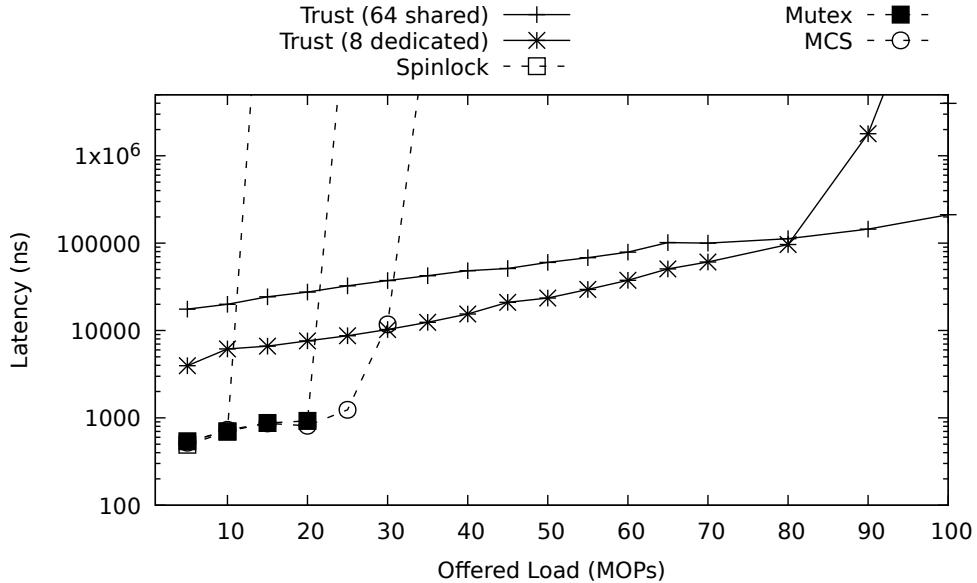


FIGURE 6. Mean latency vs. offered load. Uniform access distribution.

FIGURE 7. Mean latency vs. offered load. Zipfian access distribution,  $\alpha = 1$ .

### 6.1.2 Skewed Access Pattern: Zipfian distribution

Zipf's law [73] elegantly captures the distribution of words in written language. In brief, it says that the probability of word occurrence  $p_w$  is distributed according to the rank  $r_w$  of the word, thus:  $p_w \propto r_w^{-\alpha}$ , where  $\alpha \sim 1$ . Similar relationships, often called "power

laws”, are common in areas beyond written language [73, 63, 49, 23, 34], sometimes with a greater value for  $\alpha$ . The higher the  $\alpha$ , the more pronounced is the effect of popular keys, resulting in congestion.

Figure 5 shows the results of our fetch-and-add experiment, but with objects selected according to a Zipfian distribution ( $\alpha = 1$ ) instead of the uniform distribution above, representing a common skewed access distribution.

With this skewed access pattern, *Trust*<*T*> overwhelmingly outperforms locking across the range of table sizes. This is explained by the relatively low throughput of a single lock. In our experiments, even MCSLocks, known for their scalability, offer at best 2.5 MOPs. When a skewed access pattern concentrates accesses to a smaller number of such locks, low performance is inevitable. By comparison, a single *Trust*<*T*> trustee will reliably offer 25 MOPs, for similarly short critical sections. For more highly skewed patterns, where  $\alpha > 1$  (not shown), the curve grows ever closer to the horizontal as performance is bottlenecked by a small handful of popular items.

## 6.2 Fetch and Add: Latency

Next we measure mean latency for a scenario with 64 objects (uniform access distribution), and 1,000,000 objects (Zipfian access distribution), while varying the offered load. We show delegation results with 8 dedicated trustee cores, and with 64 shared trustee cores<sup>7</sup>. We also plot the results for a spinlock, a standard Rust mutex, and an MCS lock as above.

At low load, low contention results in low latency for locking, an ideal situation for locks. However, as load increases, the locks eventually reach capacity, resulting in a rapid rise in

---

<sup>7</sup>The evaluation system has 64 cores, 128 hardware threads. In the vast majority of cases, having both hardware threads of each core work as trustees results in reduced performance.

latency. With *Trust*< $T$ >, even low load incurs significant latency, due to message passing overhead. However, due to the much higher per-object capacity available, latency increases slowly with load until the capacity is reached. Thus, *Trust*< $T$ > offers stable performance over a wide range of loads, at the cost of increased latency at low load. The higher latency does mean that to take full advantage of delegation, applications need to have ample parallelism available.

For both Uniform and Zipfian access distributions, we also measured 99.9th percentile (tail) latency (not shown). Overall, tail latency with locking (all types) tended to be approximately 10 $\times$  the mean latency, in low-congestion settings. Delegation tail latency with a dedicated trustee, meanwhile, was 2.5 $\times$  the mean, making delegation tail latency under low load only 2–3 $\times$  that of locking.

It's also worth noting the difference between 8 dedicated trustees, and 64 trustees on threads shared with clients. The latency when sharing the thread with clients is naturally higher than when using trustees dedicated to trustee work. However, as load increases having more trustees available to share the load results in better performance. Using all the cores for trustees all the time also eliminates an important tuning knob in the system.

### 6.3 Concurrent key-value store

For a more complete end-to-end evaluation, we implement a simple TCP-based key-value store, backed by a concurrent dictionary. Here, we run a multi-threaded TCP client on one machine, and our key-value store TCP server on another, identical machine. The two machines are connected by 100 Gbps Ethernet. We compare our *Trust*< $T$ > based solution to Dashmap [69], one of the highest-performing concurrent hashmaps available as a public Rust

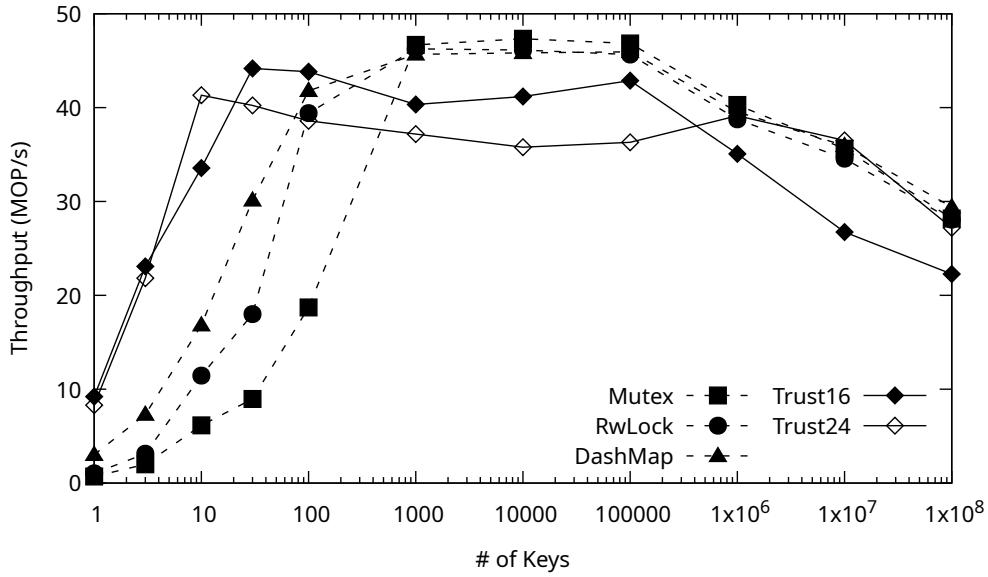


FIGURE 8. Key-value store throughput, with 5% writes and varying table size. Uniform access distribution

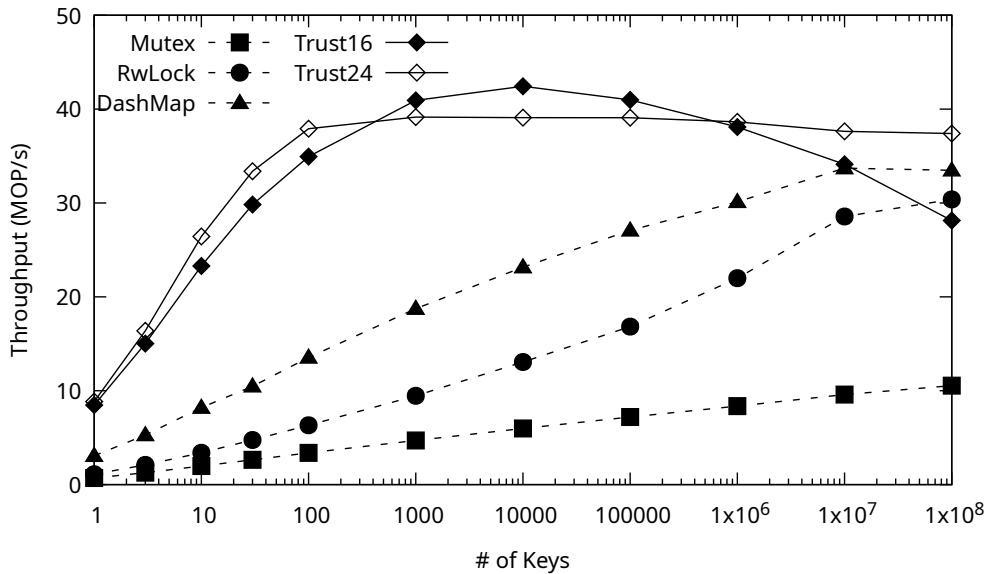


FIGURE 9. Key-value store throughput, with 5% writes and varying table size. Zipfian access distribution

crate, as well as to our own naïvely sharded Hashmap, using Mutex or Readers-writer locks and the Rust `std::collections::HashMap<K, V>`. Dashmap is a heavily optimized and

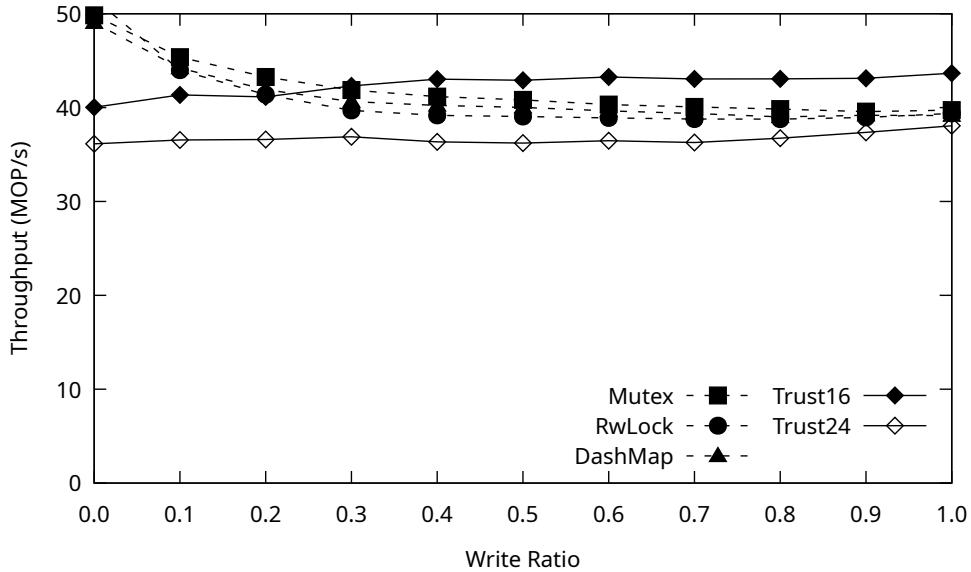


FIGURE 10. Key-value store throughput, with varying write percentage. Uniform access distribution

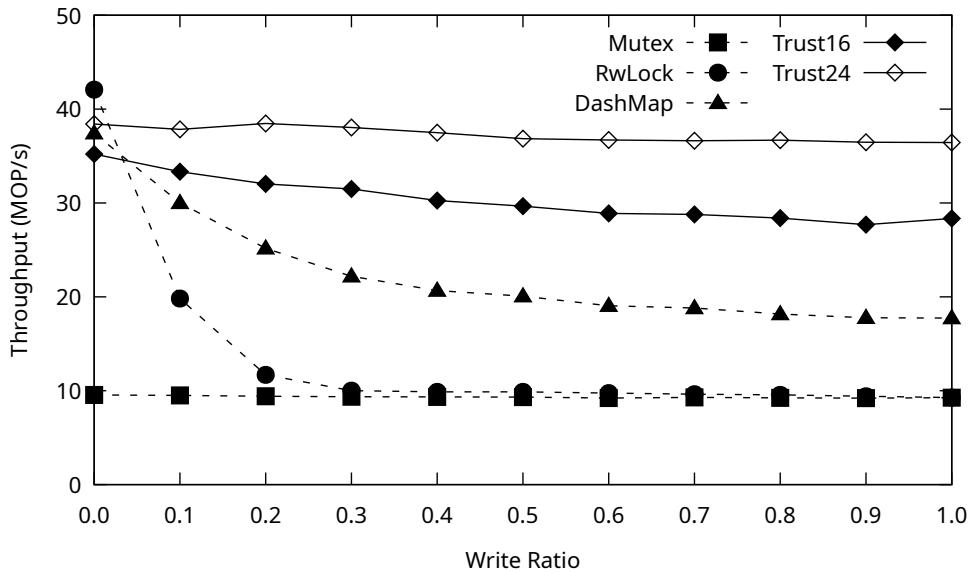


FIGURE 11. Key-value store throughput, with varying write percentage. Zipfian access distribution

well-respected hash table implementation, which is regularly benchmarked against competing designs.

We implement the key-value store as a multi-threaded server, where each worker-thread receives GET or PUT queries from one or more connections, and applies these to the back-end hashmap. Both reading requests and sending results is done in batches, so as to minimize system call overhead. Moreover, the client accepts responses out-of order, to minimize waiting. The TCP client continuously maintains a queue of parallel queries over the socket, such that the server always has new requests to serve. In the experiments, we dedicate one CPU core to each worker thread.

For our sharded hashmaps, we create a fixed set of 512 shards, using many more locks than threads to reduce lock contention. Dashmap uses sharding and readers-writer locks internally, but exposes a highly efficient concurrent hashmap API. For our *Trust*< $T$ > based key-value store, we use 16 and 24 cores to run trustees (each hosting a shard of the table) exclusively, and the remaining cores for socket workers. They are named Trust16 and Trust24, respectively. Socket workers delegate all hash table accesses to trustees. The key size is 8 bytes and the value size is 16 bytes in the experiments. Prior to each run, we pre-fill the table, and report results from an average of 10 runs.

Figure 8 and Figure 9 show the results from this small key-value store application, for a varying total number of keys with 5% write requests and 95% read request, and Uniform as well as Zipfian [73] access distributions. For Zipfian access, we use the conventional  $\alpha = 1$ . Overall, similar to the microbenchmark results, we find that the delegation-based solution performs significantly better when contention for keys is high. However, due to the considerably higher complexity of this application, the absolute numbers are lower than in our microbenchmarks. The relative advantage for delegation is also somewhat smaller, as some parts of the work of a TCP-based key-value store are already naturally parallel.

For the Uniform distribution and 5% writes, all the solutions perform similarly above 1,000 keys, a large enough number that there is no significant contention. With 100 keys and less, *Trust<T>* enjoys a large advantage even under uniform access distribution. With a Zipfian access distribution, accesses are concentrated at the higher-ranked keys, leading to congestion. In this setting, *Trust<T>* trounces the competition, offering substantially higher performance across the full 1–100,000,000 key range. It is interesting to note, also, that the Zipfian access distribution is where the carefully optimized design of Dashmap shines, while it offers a fairly limited advantage over a naïve sharded design with readers-writer locks on uniform access distributions. This speaks to the importance of efficient critical sections in the presence of lock congestion.

The throughput of Trust16 is higher than Trust24 with 1,000–100,000 keys because it is of low cost to manage a relatively small key space, while Trust16 can dedicate more resources to handle socket connections. However, the performance of Trust16 starts to degrade with more keys, because the limited number of trustees fall short when managing larger key spaces. With 24 trustees, the performance can be maintained at a high level. The difference between Trust16 and Trust24 suggests an important direction of future research: For I/O heavy processes like key-value stores, dedicated trustees will often outperform sharing the core between trustees and clients. However, it is non trivial to correctly choose the number of trustees. Automatically adjusting the number of cores dedicated to trustee work at runtime would be preferable.

In principle, readers-writer locks have a major advantage over *Trust<T>* in that they allow concurrent reader access, while *Trust<T>* exclusively allows trustees to access the

underlying data structure. To better understand this dynamic, Figure 10 and Figure 11 show key-value store throughput over a varying percentage of writes.

Here, we use 1,000 keys for the Uniform access distribution, and 10,000,000 keys for Zipfian access distribution. We note that these are table sizes where lock-based approaches hold an advantage in Figure 8 and Figure 9. For Uniform access patterns, where there is limited contention given the table size of 1,000 keys, the impact of the write percentage is muted. For lock-based designs, the performance does drop somewhat, but remains at a high level even with 100% writes.

It is interesting to note that *Trust*< $T$ > performance increases modestly with the write percentage. One reason behind this is that in our key-value store, the closures issued by reads by necessity have large return values, while the closures issued by writes have no return values at all. This may allow the trustee to use only the first, small part of the return slot, occasionally saving two LLC cache misses per round-trip.

With the Zipfian access distribution, even with 10,000,000 keys, contention remains a bigger concern, especially for Mutex. All four designs exhibit reduced performance with increased write percentages, but again, *Trust*< $T$ > proves more resilient. The efficiency advantage of Dashmap over our naïve lock-based designs is on full display with the Zipfian access distribution and a high write percentage. That said, the fundamental advantage of *Trust*< $T$ > over locking in this application is clear.

## 7 Legacy Application: Memcached

We also port memcached version 1.6.20 to *Trust*< $T$ > to demonstrate both the applicability and performance impact on legacy C applications. Memcached is a multi-threaded

key-value store application. Its primary purpose is serving PUT and GET requests with string keys and values over standard TCP/IP sockets. Internally, memcached contains a hash-table type data structure with external linkage and fine-grained per-item locking. By default, memcached is configured to use a fixed number of worker threads. Incoming connections are distributed among these worker threads. Each worker thread uses the `epoll()` system call to listen for activity on all its assigned connection. Each connection to a memcached server traverses a fairly sophisticated state machine, a pipelined design that is aimed at maximizing performance when each thread serves many concurrent connections with diverse behaviors. The state machine will process requests in this sequence: receive available incoming bytes, parse one request, process the request, enqueue the result for transmission, and transmit one or more results.

For our port to *Trust<T>*, we eliminate the use of most locks, and instead divide the internal hash table and supporting data structures into one or more shards, and delegate each shard to one of potentially multiple trustees. Thus, instead of acquiring a lock, we delegate the critical section to the appropriate trustee for the requested operation. Our ported version follows the original state machine design, with one key difference: for each incoming request on the socket, we make an asynchronous delegation request using `apply_then`, then move on to the next request without waiting for the response from the trustee. That is, rather than sequentially process each incoming request, we leverage asynchronous delegation to capture additional concurrency.

A complicating factor in this asynchronous approach results from `memcached` being initially designed for synchronous operation with locking. For any one trustee-client pair, even asynchronous delegation requests are executed in-order, and responses arrive in-order.

However, this is not guaranteed for requests issued to different trustees. Consequently, the memcached socket worker thread must order the responses before they are transmitted over the network socket to the remote client. By contrast, our delegation-native key-value store in Section 6.3 sends responses out of order over the socket, and instead includes a request ID in the response.

Another difference worth mentioning is that we don't allow delegation clients (in this case, the memcached socket worker thread) to access delegated data structures at all. This means that instead of a pointer to a value in the table, clients receive a copy of the value. This significantly improves memory locality and simplifies memory management, since every value has a single owner. However, it does incur extra copying, which may reduce performance under some circumstances.<sup>8</sup>

In practice, because memcached is written in C and *Trust<T>* is written in Rust, we cannot directly add delegation to the memcached source code. We address this in a two-step process: first, for any task that requires delegation, we create a minimal Rust function that performs that specific task. That is, a custom Rust function that becomes part of the memcached code base. Typically, such a function locates the appropriate Trust or TrusteeReference, and delegates a single closure. Second, we break out the critical sections in the C code into separate inner functions that may be called from Rust. Thus, to delegate a C critical section, we simply call the inner function from a delegated Rust closure.

Our port of Memcached to *Trust<T>* has approximately 600 lines of added, deleted or modified lines of code, out of 34,000+ lines total. This number includes approximately 200 of

---

<sup>8</sup>This can become a problem when values are large. For this use case, *Trust<T>* includes an equivalent of Rust's `Arc<T>` which allows multiple ownership of read-only values.

lines which were simply cut-and-pasted into the new inner functions for critical sections. In addition, we introduced approximately 350 new lines of Rust code, to provide the interface between the C and Rust environments.

## 7.1 Evaluation

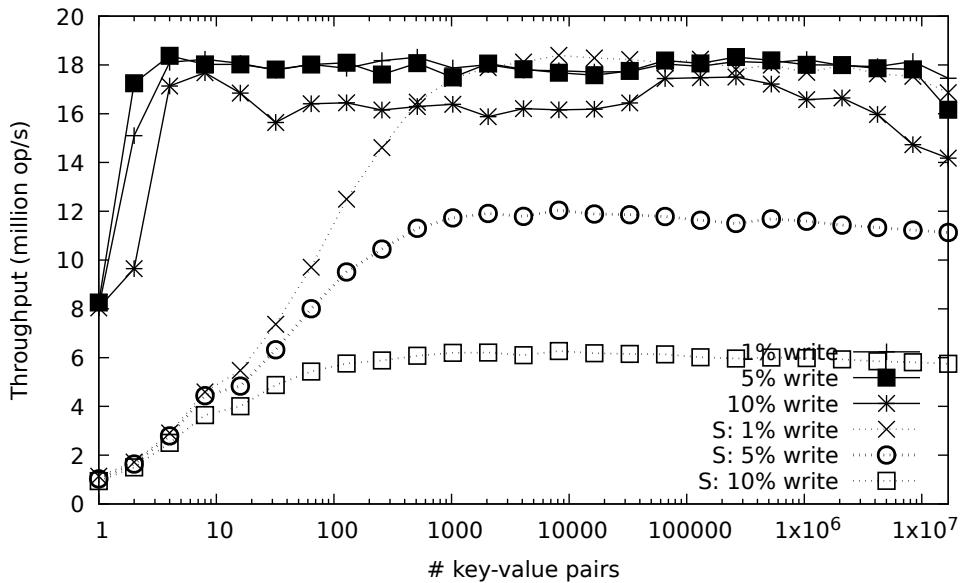


FIGURE 12. Memcached throughput with varying table size. Uniform access distribution. S: stock memcached.

To understand the performance of our delegated Memcached, we use the memtier benchmark client (version 1.4.0) with our delegated Memcached as well as stock memcached. For the cleanest results, but without loss of generality, we configure memcached with a sufficiently large hash power and available memory to eliminate table resizing and evictions. We also limit our evaluation to the conventional memcached PUT/GET operations. Recent versions of memcached feature an optional new cache eviction scheme, which trades less synchronization for the need for a separate maintenance thread. For stock memcached, we evaluated both the traditional eviction scheme and the new one. We show results for the

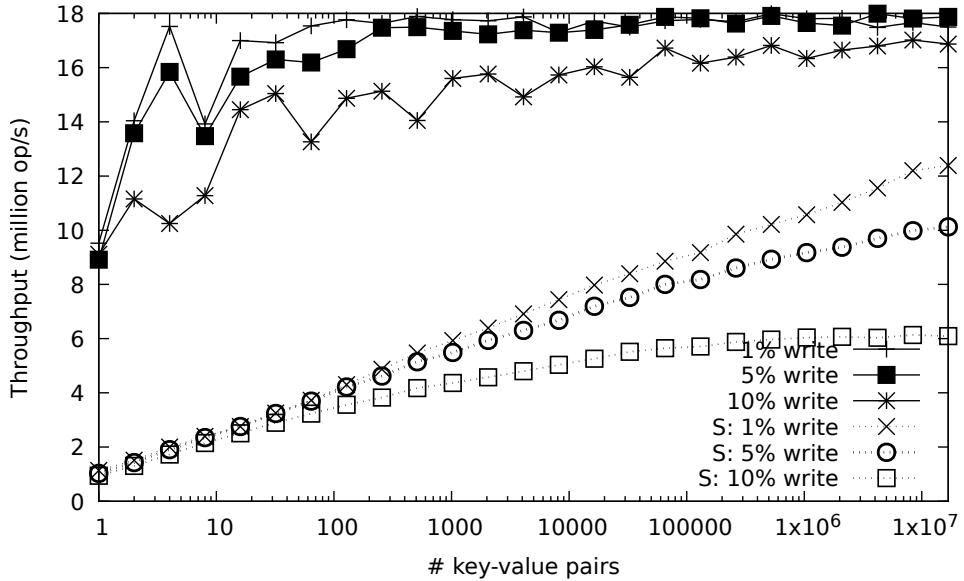


FIGURE 13. Memcached throughput with varying table size. Zipfian access distribution.

new scheme, which scales much better for write-heavy workloads and is otherwise similar in our setting. For our ported version, we use the traditional eviction scheme, maintaining one LRU per shard. Eviction is not relevant here, as we provide ample memory relative to the table size.

The server and client run on separate machines, connected by 100Gbps Mellanox-5 Ethernet interfaces via a 100Gbps switch. Both client and server machines are 28-core, two-socket systems with Intel *Sandy Bridge* CPUs and 256 GB of RAM. The machines run Ubuntu Linux with kernel version 5.15.0. Unless otherwise noted, we structure the experiments as follows: start a fresh `memcached` instance. Populate the table with the indicated number of key-value pairs, then run measurements with 1% writes, 5% writes, and 10% writes. After this, we start over with a new, empty `memcached` instance. Each data point represents a single experiment, each set to last 20 seconds. For each, unless otherwise noted, we choose `memcached` and `memtier` parameters to maximize throughput. By default, this means 28

`memcached` threads pinned to hardware threads 0–27. Running with 56 hardware threads did not yield any further performance improvement. On the memtier side, we configure 28 threads, with four clients per thread, and pipelining set to 48.

Figure 12 and Figure 13 illustrates the throughput of `memcached` as we vary the number of keys in the table. While the absolute numbers are significantly lower than in the microbenchmarks and the key-value store, the overall picture from `memcached` corresponds well with previous experiments.

Using *Trust*<*T*> results in performance improvements of more than 5× when accessing popular objects, whether this popularity is due to a uniform access distribution across a smaller number of keys, or a Zipfian distribution over millions of key-value pairs. When all items are accessed infrequently, locking suffers very little contention, and has the advantage of better distributing the work across cores. Here, this results in performance competitive with delegation, at least for read-heavy workloads.

The stock version is heavily affected by writes, due to the extra work required for these operations. This includes memory allocation, LRU updates as well as table writes, all of which involve synchronization in a lock-based design. With *Trust*<*T*>, all such operations are local to the shard/trustee, and do not require synchronization. With 5% of writes, stock memcached loses ≈40% of its performance, while the *Trust*<*T*> version sees only a minor performance penalty, resulting in delegation outperforming locking in this setting for the entire range of table sizes. While not shown, this trend continues with even more writes.

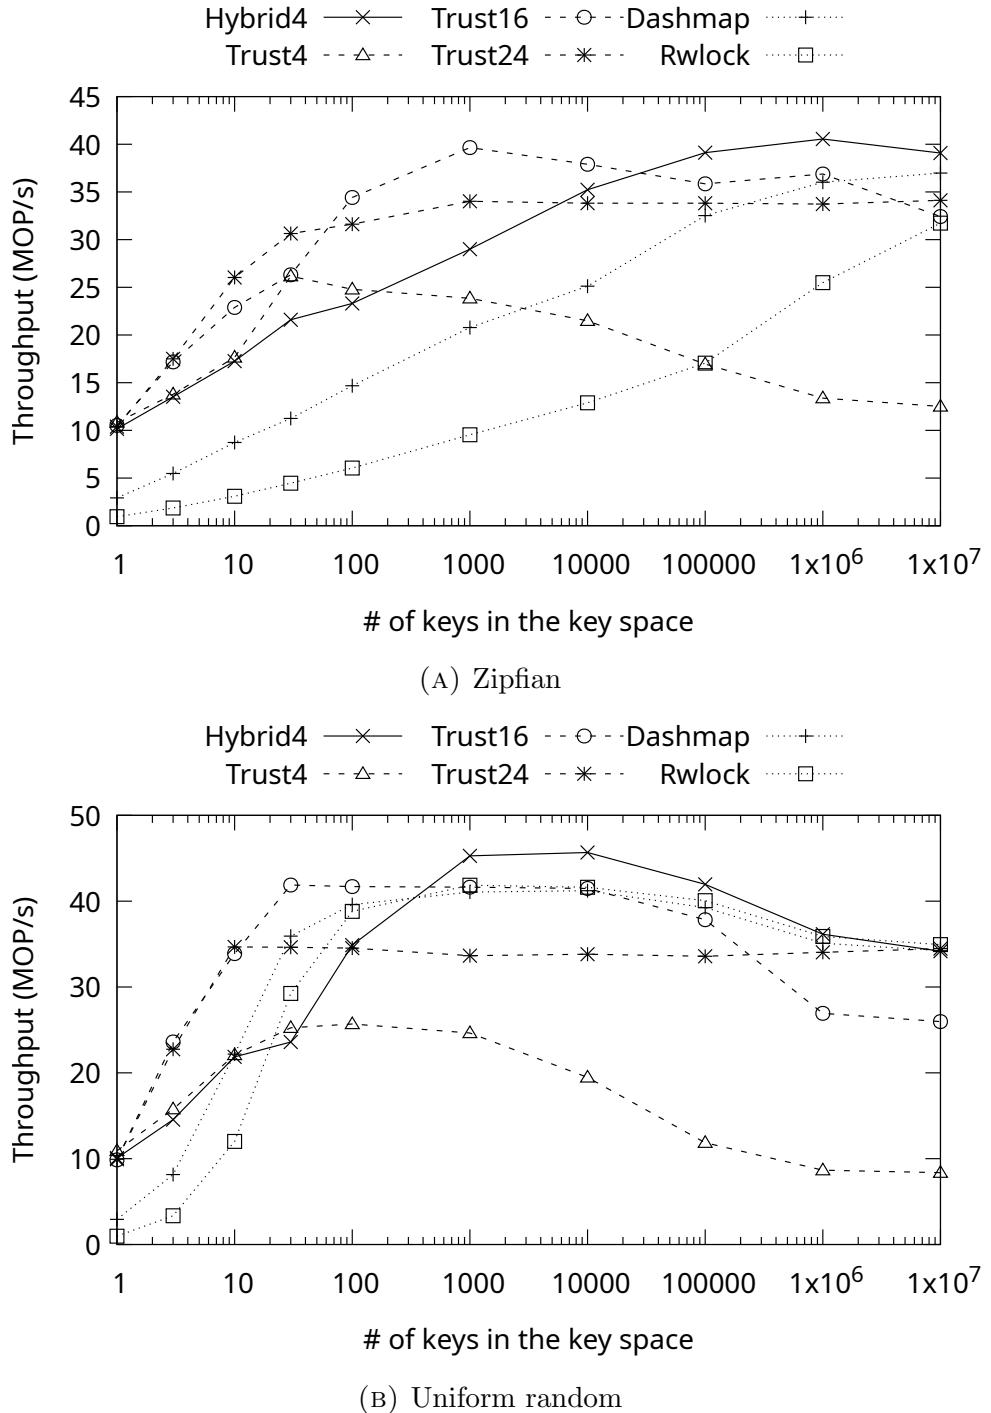


FIGURE 14. The throughput of key-value stores using compared systems with different key space size (5% write)

## 8 Hybrid Delegation

Figure 8 shows that delegation’s performance advantage depends on the workload and lock contention levels. This raises the question of whether it is possible to get the best of both worlds. In his PhD thesis, Dr. Chen [19] introduced HybridLock, a synchronization mechanism that dynamically switches between traditional locking and delegation based on workload conditions. When a HybridLock is created, it assumes low contention and starts with locking as default. If, during the runtime, it detects that the lock is highly contended, it switches to using delegation manage access to the associated property. By efficiently detecting contention and switching between the two modes, HybridLock can leverage delegation to achieve high performance in a highly contended system while keeping the overhead during low contention workload to a minimum. Figure 14 is taken from Dr. Chen’s thesis and shows that HybridLock can maintain high performance across contention levels. It is included here for reference only.

## 9 Conclusion

Synchronizing access to shared data is essential for ensuring correct behavior in multi-threaded applications. This chapter introduces *Trust< T >*, a delegation based synchronization mechanism for safe, highly performant concurrent access to shared mutable state in Rust. *Trust< T >* provides evidence that delegation can be a competitive alternative to locking in real systems. The programming model integrates cleanly in Rust, making delegation an accessible option for developers. This work lays the groundwork for additional language and runtime support to unlock the performance and scalability benefits of delegation-based designs.

## CHAPTER 4

# Gossamer

### 1 Summary

Delegation has a programming model that works very well for distributed computing. Like distributed systems, delegation does not have to rely on shared memory for operations. This makes it possible to write an application that is designed for a single machine delegated system and then adapt it for a distributed delegation system without significant changes to application code.

We introduce *Gossamer*, a distributed programming model that leverages delegation to efficiently distribute computation across machines. *Gossamer* uses Remote Direct Memory Access (RDMA) for inter-node communication, offering low latency, high throughput, and direct remote memory access without CPU involvement. This enables a unified programming abstraction that functions seamlessly on both single and multi-machine setups.

With *Gossamer*, developers can write distributed applications as if they were running on a single system. Its API allows delegation of Rust closures for remote execution, and light-weight user-space threads (*fibers*) provide efficient concurrency. Experimental results show that *Gossamer* achieves up to 3x higher throughput scaling than eRPC in microbenchmarks and matches the performance of the Graph500 reference implementation of single-source shortest path (SSSP) using MPI.

## 2 Introduction

Computing needs have grown steadily, pushing systems from simple, single-threaded execution toward far more parallel and scalable designs. Eventually, even multicore machines proved insufficient for the larger workloads, inspiring a shift toward distributed computing systems composed of many networked machines working together to handle massive data volumes, deliver high availability, and scale horizontally.

This chapter extends *Trust*< $T$ >, presented in the previous chapter, to build a rack-wide programming framework. *Trust*< $T$ > is based on delegation where a few threads (servers) act on behalf of many threads (clients). Its aim is to provide single threaded data structure performance in a multi-threaded setting. *Gossamer* takes it one step further and aims to do the same in multi-machine setting. *Gossamer* provides a very similar API to *Trust*< $T$ > enabling programmers to design applications like a normal delegation application with very few changes. These applications give the illusion of running on a single machine even though they may be distributed across many nodes. Most of the distribution of work is handled behind the scenes, but programmers have the option to customize where the worker threads should run and where data should be held in the rack if they so choose.

*Gossamer* is a rack-wide computing framework. We define rack as a set of machines connected by a single switch, also called a Top-of-Rack (ToR) switch. Our goal is to have the abstraction of a single application running across the rack. To this end we treat the rack as a single machine and if a link or a node goes down the whole rack is considered inoperable. Similarly, if a part of the application, running on any of the nodes crashes, the

whole application terminates. These aspects limit the size of the rack to a small number of machines, where we expect node failures to be very rare.

A key challenge that limits performance in a rack-wide system is the additional latency introduced by network communication. *Gossamer* mitigates this by combining RDMA with fibers. RDMA offers sub-microsecond one-way latency, significantly faster than accessing persistent storage on a single machine, though still slower than local memory and considerably slower than CPU cache accesses. Fibers help bridge this latency gap by allowing applications to issue and manage a large number of concurrent, outstanding requests. This is further amplified by batching the requests generated by all the fibers on a single system thread. While RDMA reduces latency, it does not scale efficiently as the number of connections per machine grows. To address this, *Gossamer* maintains only a single RDMA connection per hardware thread for each remote machine.

### 3 Background on RDMA

RDMA is a networking approach that allows one machine to directly access a remote machine’s memory. It uses the user-level zero-copy transfers to minimize the involvement of remote machine’s operating system and CPU as opposed to traditional TCP/IP stack that involve both on each machine heavily. There are many implementations of this concept [54, 56, 64, 53], most popular of which are InfiniBand, RoCE (RDMA over Converged Ethernet), and iWARP (internet Wide Area RDMA Protocol). The results presented in this chapter are obtained by using RoCE.

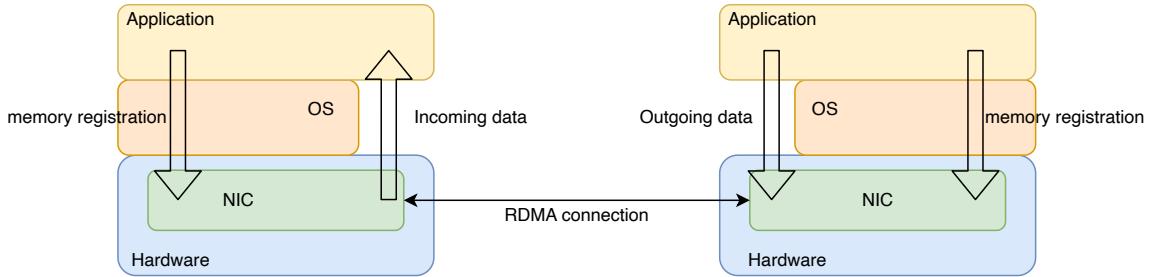


FIGURE 15. To communicate using RDMA, the applications register some memory. The NICs can then read it via DMAs as necessary.

### 3.1 RDMA API

RDMA hosts communicate using queue pairs (QPs) that consist of a send queue and a receive queue, and are maintained by the NIC. Applications post operations to these QPs by using functions called *verbs*. Each queue can be associated with a completion queue, that can be polled to learn the status of posted operations. The completion queue can be shared between both parts of the QP. For remote access the remote machine first needs to register a memory region with the NIC. The NIC driver pins this region in physical memory. The address and a key related to this region then needs to be exchanged between the machines out of band (i.e. without using RDMA). After this exchange, the remote memory can be accessed without involving either of remote operating system or CPU as shown in Figure 15. This is called *RDMA Memory Semantics*, and uses *verbs* *READ* and *WRITE*.

RDMA also provides *Messaging Semantics* that use *verbs* *SEND* and *RECV*. In this case receiver has to post a *RECV verb* before the sender can send the data. In this regard it is similar to an unbuffered socket's implementation. Just like *Memory Semantics*, *Messaging Semantics* also bypass the remote kernel but unlike *Memory Semantics*, it has to involve remote CPU to post a *RECV*. These *verbs* also have slightly lower latency than *READ* and *WRITE* [46, 25].

### 3.2 Transport types

RDMA transports are either connected or unconnected (also called datagram), and either reliable or unreliable. Connected transport require a one-to-one connection between two QPs. If an application wants to communicate with  $N$  machines, it will need to create  $N$  QPs. With unconnected transport one QP can communicate with many QPs. For reliable transport NIC uses acknowledgments to guarantee in-order delivery and return an error code on failure, while unreliable transport does not provide any such guarantee. InfiniBand and RoCE use lossless link layer, so even in case of unreliable transports, losses are pretty rare and happen because of bit error or link failure. In case of connected transport a failure will break the connection. Not all transports provide all of the verbs. Table 5 gives an overview of transport types and the verbs they support. Current implementations of RDMA only provide reliable connected (RC), unreliable connected (UC) and unreliable datagram.

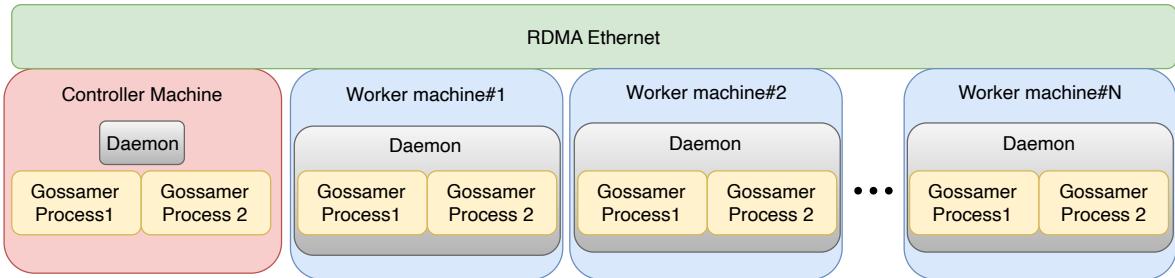
The meaning of a success status while polling a completion queue depends on the type of transport the related send/receive queues are using. A work completion is generated for all receive requests and added to the completion queue. For send requests the work completion is only generated if the signaled bit was set to true in the work request or the send request ended with an error. In case of reliable transport, a success for send requests indicates that the data has been written to the remote machine's memory. For unreliable transport however, a success only means that the local buffers can be reused safely, with no indication of the remote memory's status. A successful work completion also indicates that any unsignaled work posted earlier has finished successfully [11]. *Gossamer* uses RC QPs to ensure any data loss in the network is caught. We also use the cumulative nature of work

Verb	RC	UC	UD
SEND/RECV	✓	✓	✓
WRITE	✓	✓	✗
READ	✓	✗	✗

TABLE 5. Operations supported by each transport type.

completion successes to our advantage by not using signaled send for every RDMA write and only having a signaled send after 1000 unsignaled sends. This allows us to cut down on time spent while polling the completion queue.

## 4 Gossamer Design

FIGURE 16. High-level design of *Gossamer*

*Gossamer* allows developers to write rack-wide applications just like single machine applications. Programmers will use the delegation framework to write applications that are distributed by *Gossamer* with very few changes to the code. Figure 16 shows a high level view of what a rack with multiple *Gossamer* applications looks like. Developers use the controller machine to interact with the rack and launch applications. Here controller machine is just a normal machine running *Gossamer* processes like any other machine in the rack, the only significance is that this is the machine that users log into for terminal access. Each machine in the rack has a daemon running that is always listening for instructions from the controller machine. The daemon is responsible for determining the topology of the rack and

managing the applications. A machine can have multiple applications running on it at the same time. Rack applications that are running simultaneously are isolated from each other like multiple processes on a single machine. The rest of this section describes some design details that are particularly interesting.

#### 4.1 Gossamer Daemon

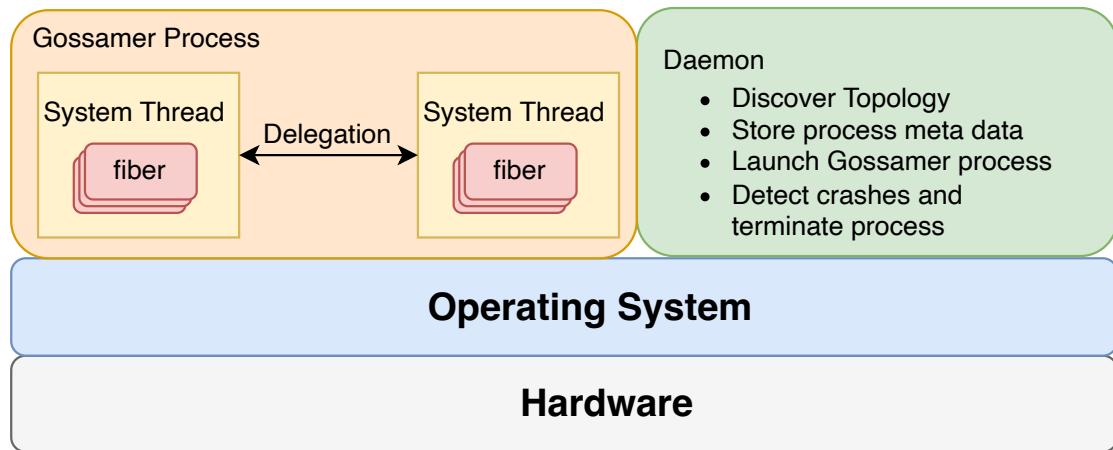


FIGURE 17. A *Gossamer* process on one machine

As shown in Figure 16, each machine has *Gossamer* daemon running. These daemons connect with each other to discover the rack size and store metadata like gossamer-id, machine-ip etc. The daemon is responsible for managing any *Gossamer* processes running on the rack at any time. To launch a *Gossamer* application, users need to run the application binary on the controller machine only. The application will contact the local daemon and tell it to start the process specifying how many system threads the whole *Gossamer* process should use. This daemon will contact the rest of the rack and tell the daemons there to launch the process as children. The daemon is also responsible for detecting crashes on any local instance of a *Gossamer* process and terminating it across the whole rack. This means that

*Gossamer* processes are fate sharing in the same way single machine processes crash if a single thread crashes.

## 4.2 Gossamer Process

To launch a *Gossamer* process, the user logs into the controller machine and starts the application there. Since parts of *Gossamer* rely on similar memory layout (as discussed in Section 4.6), each machine in the rack should be populated with the same binary ahead of time. The application needs to have its *main* function call a helper function from the *Gossamer* library and pass a function that would act as the real *main* function along with how many threads the application needs to use across the whole rack. This function will be referred to as *gsm\_main* for the rest of the chapter. The process consists of many fibers per system thread for concurrent operation as shown in Figure 17. Fibers are lightweight, user space threads that can spawn on any system thread that is part of the *Gossamer* process, and on any machine in the rack. Fibers that need to access shared data make delegation requests consisting of closures that modify the data as needed. From the point of view of the operating system, a single instance of *Gossamer* process behaves like any other normal process. The main difference is that instead of using system calls like *getpid()*, a *Gossamer* process contacts the daemon to get any metadata needed.

## 4.3 Trustee, Trust and Property

*Gossamer* builds on the Trust/Trustee concept from [7] (Chapter 3 in this dissertation), which we introduce briefly here. A *Trustee* is a worker fiber that processes delegation requests and sends back the response. The application can entrust some data to a trustee if the delegated work needs access to it and receive a *Trust* that holds relevant metadata. This

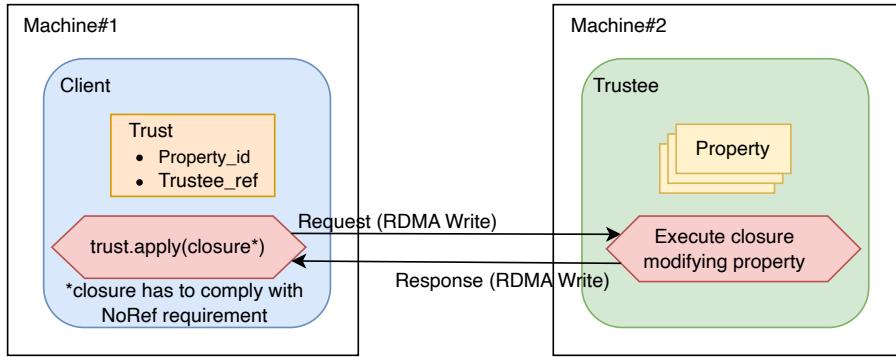


FIGURE 18. Trustee is a worker and clients can delegate work using a trust that holds information about property and trustee

```

1 pub auto trait NoRef {}
2 impl<T> !NoRef for &T {}
3 impl<T> !NoRef for *T {}
4 fn apply<F:Closure+NoRef>(closure: F);

```

TABLE 6. Auto traits combined with negative trait implementations allow functions to restrict closures that capture pointers.

entrusted data is referred to as *Property*. The Trust can be used to delegate any work that needs access to this property. Figure 18 shows the relationship between Trust, Trustee and Property. A Trust can be cloned and shared with other fibers so that they can also start delegating work that needs access to same property. Multiple properties can be entrusted to a single trustee at the same time. Trustees can act as remote or local trustees based on where the fiber trying to access the property lives in the rack.

#### 4.4 Pointers and Serialized Data

*Gossamer* leverages the Rust type system to prevent the sending of references over delegation channels. One of the first challenges that arise when extending delegation to multiple machines is the fact that any pointers or references captured by a delegated closure will not be valid on a remote machine. This reference or pointer will result in undefined behavior on

```

1 // trustee_ref can be a reference to either local or remote trustee
2 let trust = trustee_ref.entrust(0u64);
3 let data = Box::new(5);
4 trust.apply(move |property| *property + *data);

```

(A) If the delegated closure captures a pointer, it can result in a segmentation fault.

```

1 // trustee_ref can be a reference to either local or remote trustee
2 let trust = trustee_ref.entrust(0u64);
3 let data = Box::new(5);
4 trust.apply_with(move |property, data| {*property + *data}, data);

```

(B) Instead of capturing a pointer, the relevant data can be serialized and sent to trustee.

TABLE 7. *Gossamer* disallows pointers from being captured and uses serialization to get around the limitation.

```

Compiling trust_eval_rdma v0.1.0 (/home/nahmad30/remote-delegation/trust_eval_rdma)
error[E0277]: the trait bound `*const u64: NoRef` is not satisfied in `'{closure@src/main.rs:70:17: 70:32}'`
--> src/main.rs:70:17
70 |     trust.apply(move |property| *property + *data);
   |     ^^^^^^
   |
   |     *const u64 contains references or pointers, and cannot be passed in closures or return values on
   |     the delegation channel. First, make sure you are passing a move closure. If you already did that, try passing values
   |     of type *const u64 in the withparameter of apply_with instead: apply_with(|property, ref_obj| ..., ref_obj). You can p
   |     ass multiple values in the with parameter by enclosing them in a tuple:apply_with(|property, (o1,o2)| ..., (ref_obj1,
   |     ref_obj2))
   |     |     within this `'{closure@src/main.rs:70:17: 70:32}'`
   |     |     required by a bound introduced by this call
   |
   |= help: within `'{closure@src/main.rs:70:17: 70:32}'`, the trait `NoRef` is not implemented for `*const u64`, which
is required by `'{closure@src/main.rs:70:17: 70:32}: NoRef`
```

FIGURE 19. *NoRef* trait forces a compile-time error message instead of runtime undefined behavior, if pointers are captured by a delegated closure.

any machine other than where it originated. To prevent this, *Gossamer* uses a combination of rust features named *auto\_traits* and *negative\_impls*. In rust, traits define the behavior of data types. They inform the compiler what functionality a type has and can be used to restrict which data types can be passed to a generic function. *Auto\_traits* is a feature that automatically implements a trait for every data type, be it an existing type or user defined. A negative implementation is used to exclude a data type from the auto implementation. *Gossamer* defines an auto trait called *NoRef*. Then all the types that contain a reference

or raw pointer are given negative implementation. The function *apply*, that is used to send delegation messages as described in previous section, requires all the closures to comply with the NoRef trait as shown in Figure 18. This is implemented using the code shown in Table 6. This however limits severely what type of data can be captured by any delegated closure as many frequently used types like strings and vectors contain pointers internally. To get around that *Gossamer* provides a way to send any data type that can be serialized along with the closure as an extra parameter to the delegated closure. Table 7 shows a code example where a Box, a Rust smart pointer, is serialized instead of being captured. If the programmer makes a mistake and tries to use a closure that has captured a reference for delegation, a custom error message at compile time will inform them what the issue is and direct them to use the serialization method instead. An example of this message is shown in Figure 19.

## 4.5 Request size and TCP fallback

The requests going from a client thread to the same trustee are batched and on trustee side the requests coming from a single thread are all contiguous in memory. This is achieved by the trustee having a pre-allocated space in memory (called request slot) for incoming requests from each client thread. This limits how many requests a client thread can send in a single batch depending on how much data is being sent with each request. As an example if each request captures 32 bytes of captured data, the total request size is 40 bytes. Next section goes into detail about the format of a *Gossamer* delegation request. Assuming 1 KB request slot, one batch can have at most 25 requests. While this allows us to optimize for small requests it also places a hard limit on how much data a request can capture i.e. the size

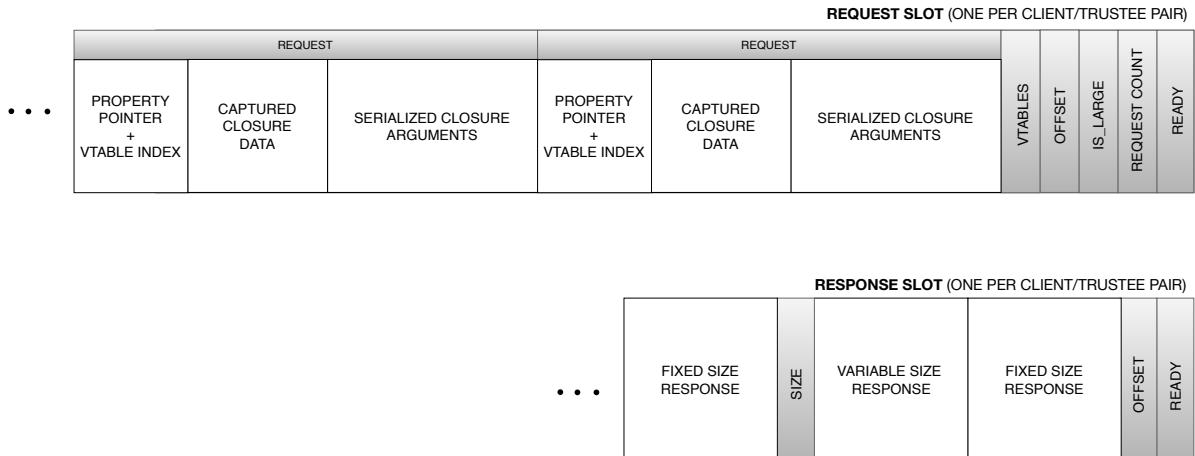


FIGURE 20. *Gossamer* Request and Response slot layout

of the request slot. To work around this limit and support larger requests, *Gossamer* uses TCP to send any captured data larger than that separately, while the request header goes in the request slot as normal. This degrades the performance severely as TCP is much slower than RDMA. We use TCP here instead of RDMA because, as mentioned in Section 3.1, RDMA needs the memory used to be pre-registered with the NIC, placing an upper limit on how much data can be sent using RDMA at once, i.e. the size of pre-registered buffer. Since there would be an upper limit anyway and larger requests are not the focus for *Gossamer*, we decided to opt for the simpler design in place at the moment.

#### 4.6 Request/Response Slot Format

As discussed in the previous section, the number of requests sent in a batch depends on the size of the requests plus the size of any headers sent with the request. Section 5.3 describes the format of request and response slots used for single machine delegation channel. Here we describe further optimizations made to request slots to minimize per request metadata to reduce wasted network bandwidth. The two-slot optimization of single machine

delegation channel is also disabled for remote delegation, as the cost of doing two separate RDMA write operations outweighs any benefit gained from it.

The necessary parts to process a *Gossamer* request are as follows:

- The closure. Closures in rust are fat pointers consisting of a vtable that points to where the code is in the text section along with the size of captured data, and a data pointer that points to the captured data.
- The pointer to where the property is located on trustee.
- Length of serialized data.
- The captured data.
- Any serialized data.

While we can't avoid sending the captured data, serialized data and the property pointer, *Gossamer* employs some strategies to minimize the information that needs to be sent inside the request slot. As the data pointer within the closure is not going to be valid on a remote machine, there is no need to include that in the request. Since the same binary is running on all the machines in the system, they all have access to any information known at compile time. This includes all of the information in the vtable, provided the vtable pointer. Here, we make the observation that for the vast majority of the runtime of an application, most of the requests in the request slot will be about a single closure or a few closures at most, meaning we can have a small lookup table for a few vtable pointers in the request slot instead of including one with every request. Combining that with the start point of data received, the trustee can reconstitute the closure locally, removing the need to send the closure in

the request slot. If there are indeed more closures in the request slot than will fit in the lookup table, any extras can be sent in the request slot. This will reduce the data being sent in the common case with minimal overhead, but will be no more expensive in the special case. Similarly, if the serialized data has a size known at compile time the trustee already has access to it, so the only time it needs to be sent with the request is if it is not known at compile time. One thing to note here is that we rely on same vtable pointer to be the valid on all of the machines which is not the case normally due to linux address space layout randomization (ASLR), making disabling it a necessity for *Gossamer* to function.

Similarly, the responses for all of the requests in a request slot are batched in a response slot and written to the client's memory with a single RDMA write. Unlike a request however, the size of a response is not always known when submitting a request. If the response size cannot be determined statically, it is preceded by 8 bytes that contain its size. While a known response size can be used to limit the number of requests in a request slot just like the size of requests, unknown size responses can lead to scenarios where not all of the responses for a request slot fit in a single response slot. In such cases the trustee will send anything that doesn't fit in the response slot to the client using TCP in a similar fashion as discussed in Section 4.5. In addition, if the response size is zero and statically known, then nothing is sent in the slot for that particular response. Figure 20 shows the format for both the request and response slots.

In addition to the request/response, the slots also contain some metadata. They both contain a ready bit and an offset, with the request slot containing a few more things. The

ready bit is there to let the other side know of a new request or response. It has to be at the end to ensure that when a client or trustee sees the ready bit the rest of the slot has already been written. This is necessary because RDMA writes data sequentially and having the start of a slot been written to, does not mean all of it is available to read. The offset has to do with another optimization for the size of RDMA writes. Depending on the size of requests and responses, a slot can be sent to the remote machine before it is 100% full, in which case writing the full slot to remote machine's memory wastes bandwidth and lowers performance. Since the ready bit needs to be at the end of the slot and in a known place, requests and responses are aligned to the end of the slot leaving the start of it empty. The offset informs the remote side where the actual data starts in the slot. The request slot also contains the request count and the vtable cache as discussed earlier. The last thing in request slot metadata is a bit that indicates if there is a single large request in the slot or multiple small ones. As discussed in Section 4.5, if the request will not fit in the slot it is sent using TCP and only the relevant pointers go in the slot. The decision to only send a single request if it is large is to simplify the design since large requests are not the focus of this work.

## 4.7 Trustee and Client memory layout

The clients and trustees on a single machine share memory, so a client can prepare the whole request slot in a pre-designated place and then flip the ready bit at the end to signal the trustee that it is ready. The trustee continuously polls this piece of memory to determine the arrival of new requests. This leads to  $\text{num\_threads}^2 \times \text{sizeof}(\text{request\_slot})$  bytes of memory being reserved for request slots. Same for response slots as well. We can think

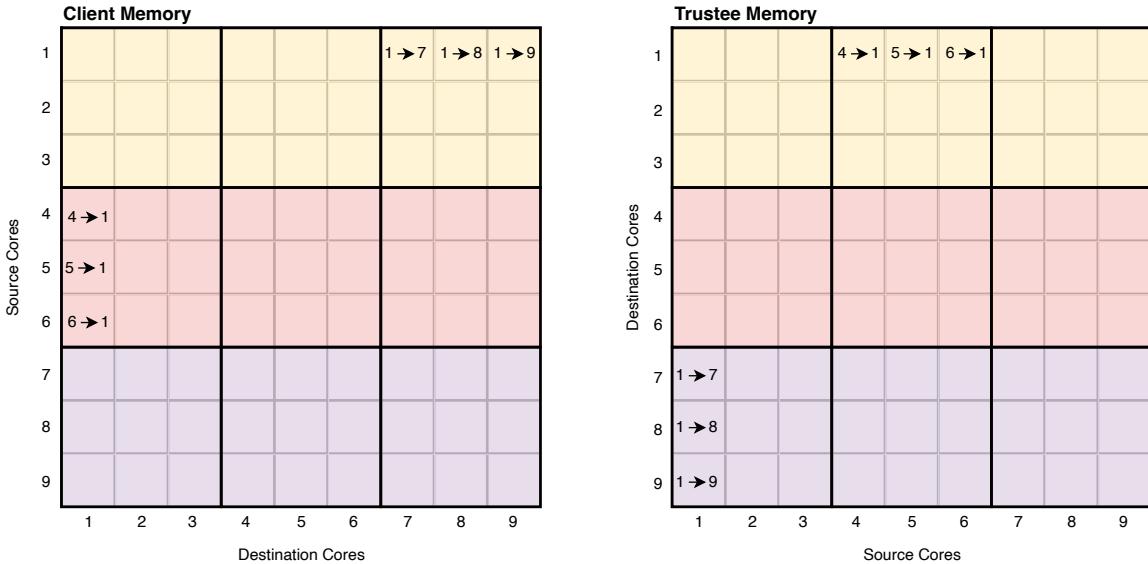


FIGURE 21. Memory layout for Request Slots in client and trustee memory for a rack of three machines where each machine has three system threads.

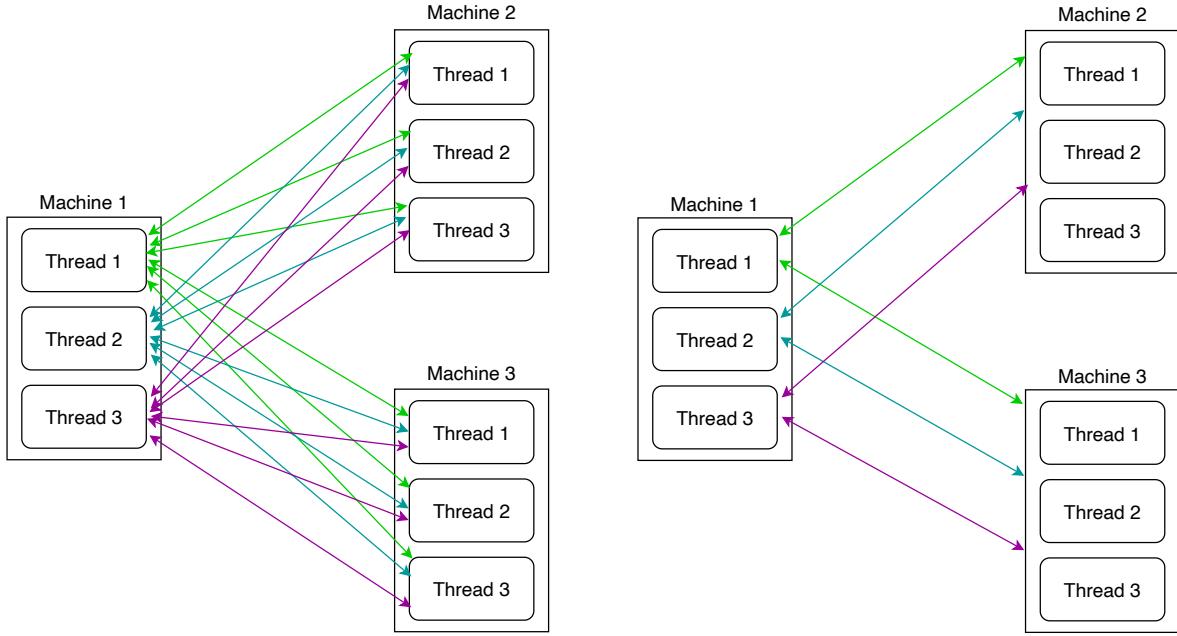
of that as a *square* in memory with `num_threads` side length. When we move on to a distributed setting, we can extend this *square* to have `num_threads × num_machines` side length. The whole of the bigger *square* does not need to be allocated on each machine as that would waste memory. Instead, parts of it are allocated on each machine. Figure 21 shows an example layout for a system with three machines that have three system threads each for a total of nine threads. In the figure, each cell represents a request slot going from a source thread to a destination thread as labeled and horizontal rows represent memory that is physically contiguous while different colors divide the *square* based on which machine the memory is allocated on. Here, unlike in a single machine setting, there are two *squares*. Clients on one machine need to prepare the requests in local memory before writing the request slot to the trustee’s memory on another machine. Theoretically, RDMA allows the clients to write the requests in the trustee’s memory as they come and flip the ready bit once

the slot is full similar to single machine channel, however this would result in many small RDMA write operations, reducing the effectiveness of batching and sinking the performance.

From the perspective of the whole rack as a single system, each row of request slots in the client memory *square* maps to a column in the trustee memory *square*. For example a request slot going from thread 1 (thread 1, machine 1) to thread 7 (thread 1 machine 3) is prepared in memory represented by the cell in row one and column seven in client memory *square* and then written to the cell in row seven and column one in trustee memory *square* and so on. Figure 21 also shows a few additional examples. There is a similar system in place for response slots as well.

#### 4.8 Scaling impact of increasing number of Queue Pairs

As discussed in Section 3.1, RDMA communications happen via a queue pair (QP). Theoretically each machine pair only needs to have a single QP resulting in `num_machines`<sup>2</sup> QPs. However, this means sharing a QP among many threads. The QPs use a spin lock internally to make sure QPs can be shared safely, but that results in threads having to wait for each other before they can send requests to a trustee. A better but naive design might be to have a QP for every pair of system threads in the whole rack resulting in QPs in the order of  $(\text{num\_machines} \times \text{num\_threads\_per\_machine})^2$ . This also results in lost performance due to NIC cache limitation. The NIC caches any recently used QPs, along with page mappings for any recently used pinned pages in its SRAM. Increasing the number of QPs results in NIC running out of SRAM, which in turn results in cache misses. We use a middle of the road system where instead of each thread making a separate QP for every other thread, it makes a QP for each machine in the rack. This reduces the total number



(A) An RDMA connection between each pair of threads means many connections per machine, resulting in lost performance.  
 (B) Having client threads use the same connections for all trustees on a remote machine avoids performance loss.

FIGURE 22. Quadratic vs Reduced connections between threads.

of QPs in the system to  $\text{num\_machines}^2 \times \text{num\_threads\_per\_machine}$ . Figure 22 shows the reduction in the number of connections gained by using this design. Figure 23 shows an experiment that demonstrates the performance gained by reducing the number of QPs in the system. This experiment was performed with 6 machines, all connected to switch with 100Gb/s ethernet links. Each thread writes 256 bytes in a randomly chosen remote thread's memory twenty million times. As the number of threads increases, thereby increasing the number of QPs, the system with reduced number of QPs maintains peak throughput whereas the system with quadratic number of QPs starts to fall off, resulting in a 25% decrease at the extreme end.

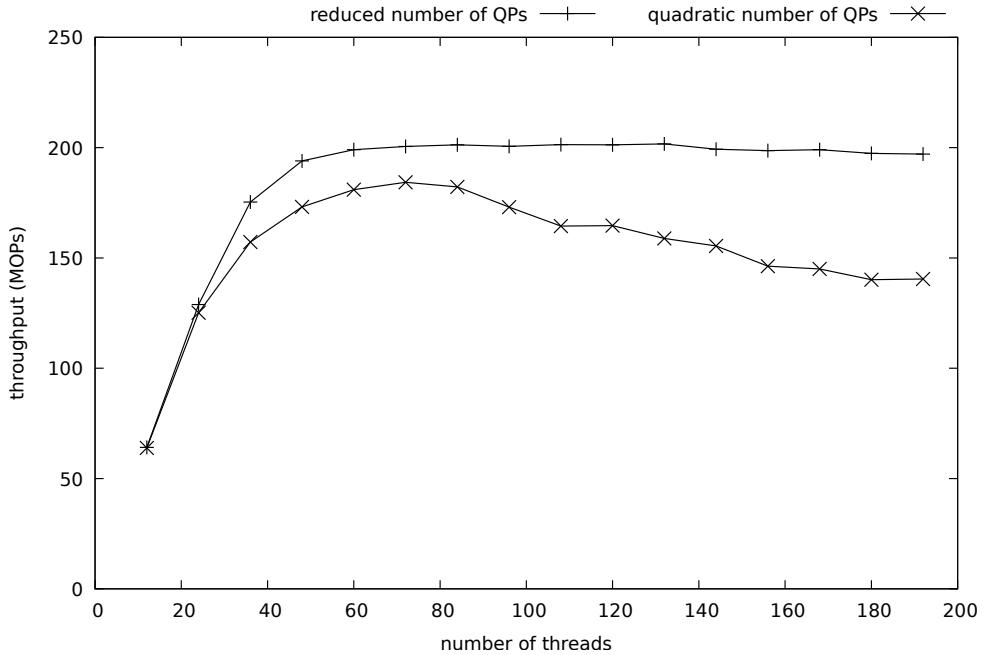


FIGURE 23. Throughput of RDMA writes in Millions of operations per second (MOPs) for quadratic and reduced number of QPs. Each write is 256 bytes.

## 5 Evaluation

In this section we use microbenchmarks to evaluate *Gossamer* on throughput and latency.

We also use two application benchmarks to evaluate how *Gossamer* performs in the real world. These experiments were performed on the r6615 cluster from Cloudlab Clemson using 12 machines. Each machine here uses AMD EPYC 9354P CPU and has 32 threads over four NUMA nodes at 3.3 GHz. We do not use hyperthreading for this experiment. Each machine has 100 Gbps Mellanox ConnectX-6 interfaces and connected to a Dell Z9432 switch. These machines are equipped with 188 GB RAM each and are running Ubuntu 22.04.2 LTS (GNU/Linux 5.15.0-160-generic x86\_64).

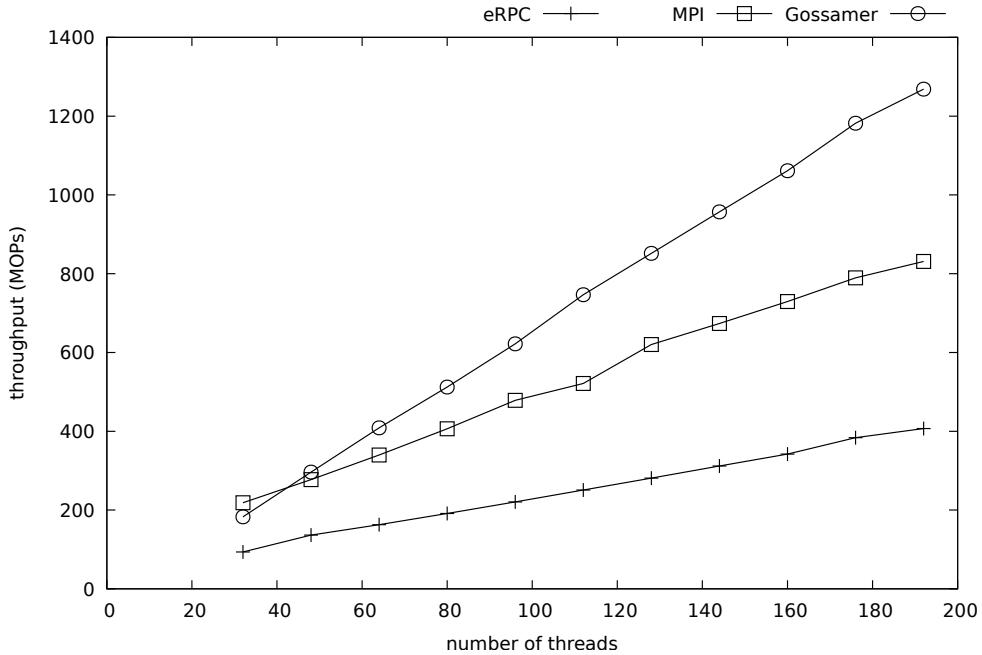


FIGURE 24. Throughput across eRPC, MPI, and *Gossamer* in an Echo application. Each request here is 32 bytes.

### 5.1 Microbenchmark: Echo Throughput

We use an echo application to evaluate the throughput scaling of *Gossamer* as the number of machines in the rack increases. For this experiment, each trustee/server sends the same data it received as a response. Each thread acts as both client and server. Clients choose a server at random from a uniform distribution and send a request containing 32 bytes of data. Each client sends ten million asynchronous echo requests. We compare against eRPC [35] and an MPI library used by and developed for graph500 reference benchmark [3].

Figure 24 shows the results of our experiment. Here we use 16 threads per machine as the eRPC benchmark only supports up to two NUMA nodes. Each successive data point has one more machine than the previous one, resulting in an increment of 16 threads. eRPC starts with 2.9 MOPs per thread and maintains that as more threads are added. This is in line with their reported performance in a github issue [2] for two machines connected with a

100 Gbps connection. *Gossamer* starts with double the throughput and scales 3x better than eRPC. Similarly, though the MPI application starts out performing better than *Gossamer*, it falls behind quickly as the number of machines increases.

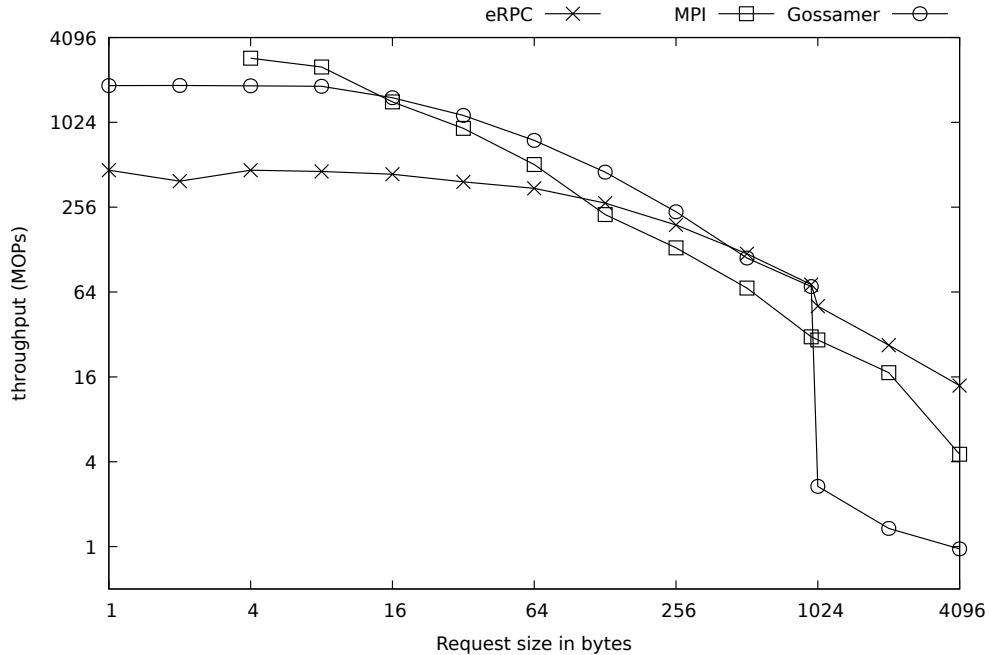


FIGURE 25. Throughput across a range of request sizes for eRPC, MPI, and *Gossamer* in an Echo application. This experiment uses 192 threads across 12 machines.

Figure 25 shows the performance of all three systems across a range of request sizes. The MPI system performs better than both *Gossamer* and eRPC due to smaller request headers, however as the request size increases, this advantage becomes negligible and the other systems overtake it. All three systems experience a drop in the number of requests served as the request size increases, however *Gossamer* experiences a particularly steep drop as it crosses the 1KB threshold. This is explained by the TCP fallback for large requests as discussed in Section 4.5.

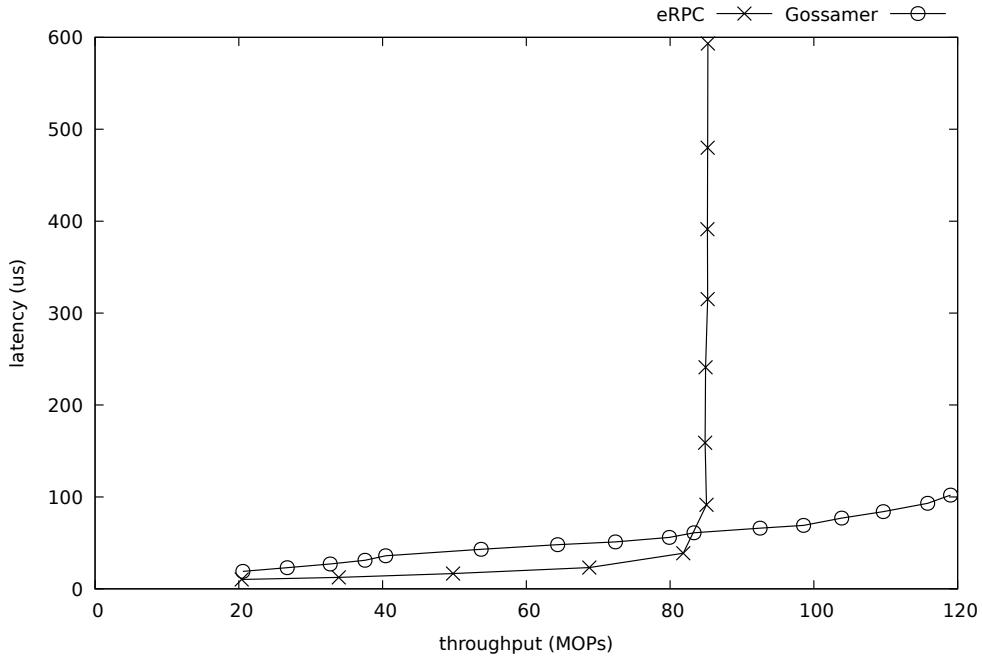


FIGURE 26. Throughput vs latency under load for *Gossamer* and eRPC in an Echo application with two machines and 16 threads per machine.

## 5.2 Microbenchmark: Echo Latency

Next we evaluate 99.9th percentile latency under load for eRPC and *Gossamer*. We exclude the MPI application from this experiment as the MPI library we use is highly optimized for throughput at the cost of latency, resulting in it having minimum latency of 1 millisecond. This experiment also uses the echo application with a request size of 32 bytes. As both *Gossamer* and eRPC use batching for asynchronous requests, load here is represented by the number of outstanding requests at any given time.

Figure 26 shows the results of this experiment. Each data point here represents a doubling of batch size until 32 and then an increment of 16. Both system have increased throughput and latency as the number of outstanding requests increases, however eRPC does better at smaller loads in terms of both latency and throughput. This is because of the optimizations discussed in Section 4.6, that result in better performance at higher load at the cost of

small performance penalty at lower loads, along with the overhead of switching fibers being significant at lower load. eRPC then quickly reaches its peak throughput and any increase in load only increases the latency whereas *Gossamer* is able to achieve much higher throughput while remaining below 100  $\mu$ s tail latency.

### 5.3 Application: Key-Value Store

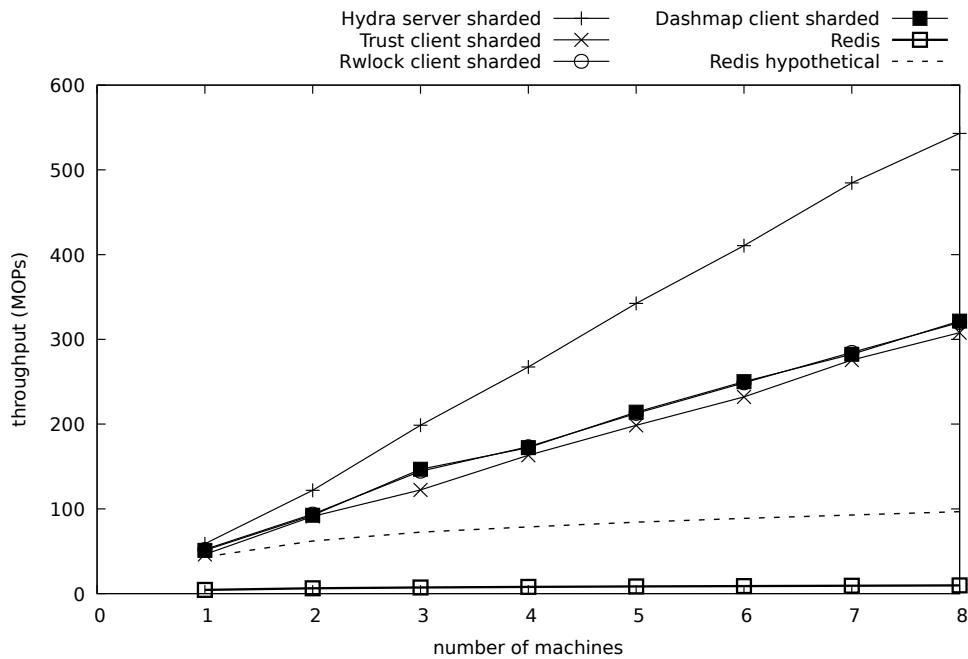


FIGURE 27. Throughput of key-value stores on up to 8 server clusters and 100% read workload with uniform distribution.

To understand the real world performance of *Gossamer*, we implemented a distributed key-value store named Hydra. Hydra divides the key space among all of the server machines equally and each server in turn divides the relevant key space among a set of dedicated trustees. A client can connect with any server that is part of the store and submit queries. Hydra uses *Gossamer* to respond the queries using delegation. We compare Hydra with Redis Cluster [4], a common production ready key-value store. We also compare against a

## Chapter 4. Gossamer

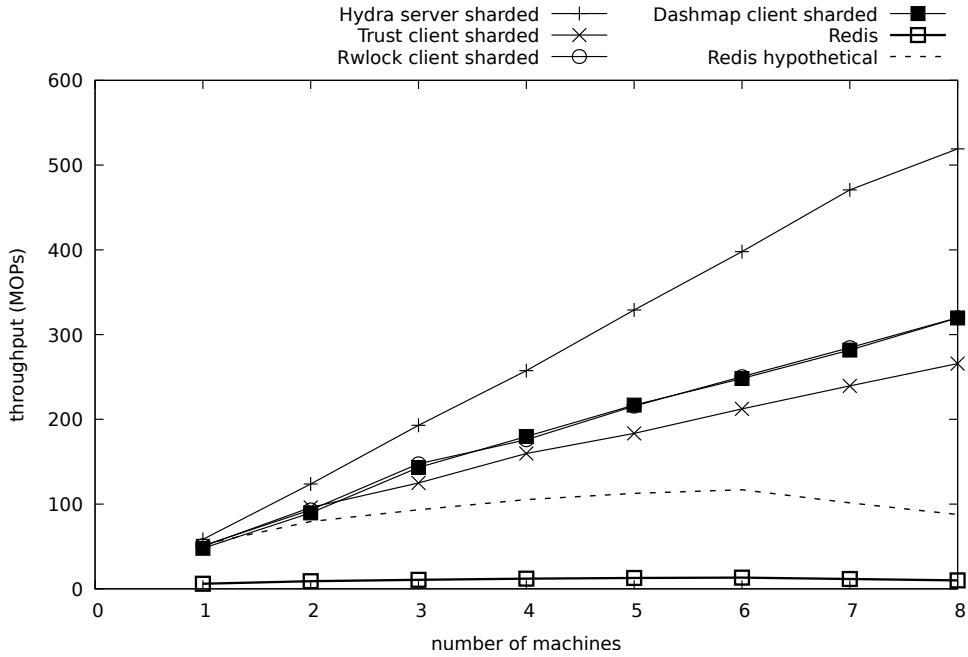


FIGURE 28. Throughput of key-value stores on up to 8 server clusters and 5% write workload with uniform distribution.

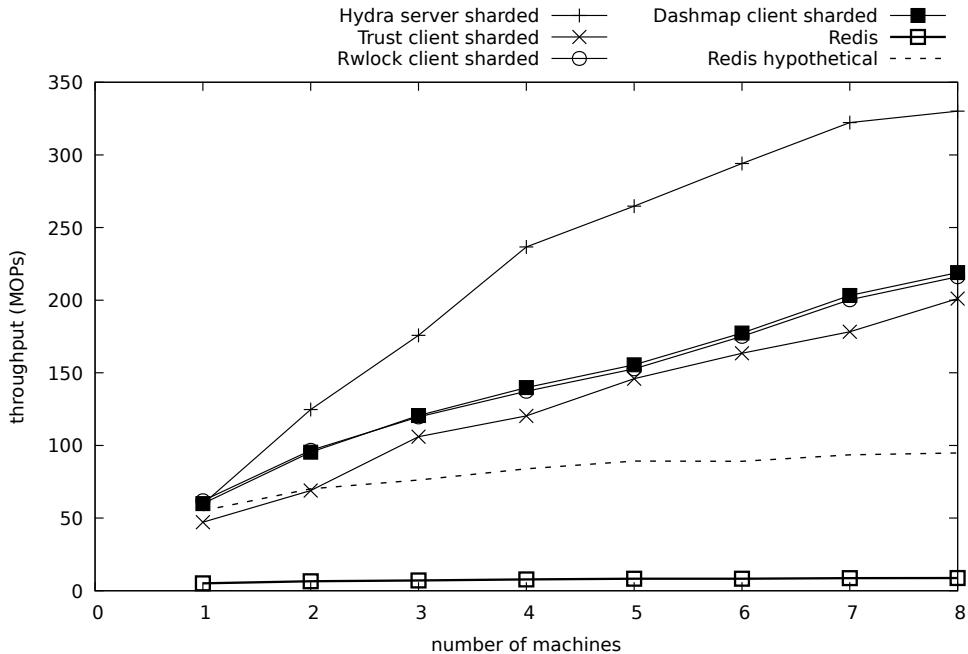


FIGURE 29. Throughput of key-value stores on up to 8 server clusters and 100% read workload with Zipfian distribution.

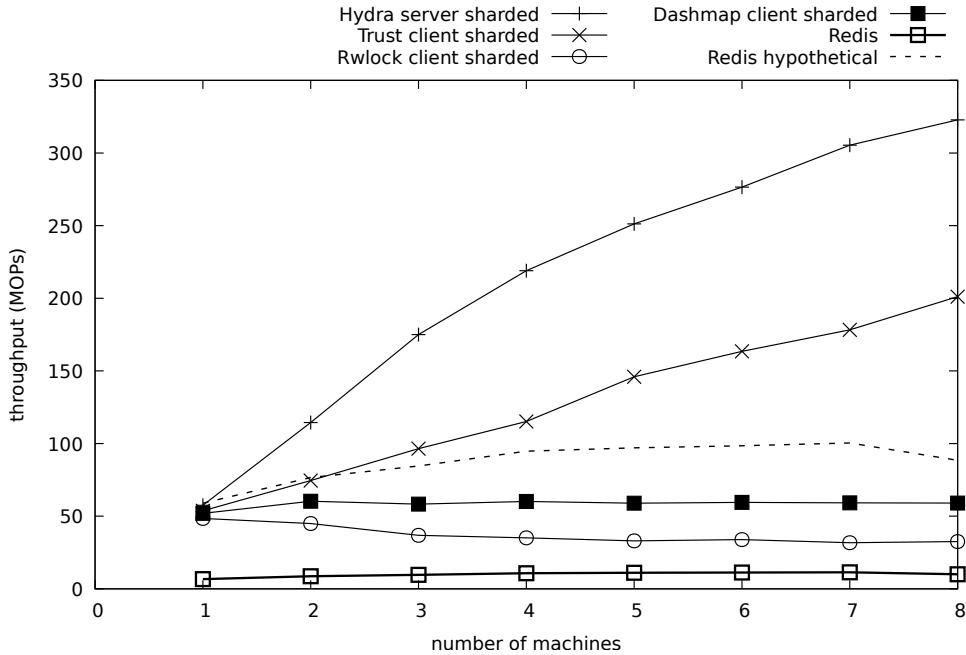


FIGURE 30. Throughput of key-value stores on up to 8 server clusters and 5% write workload with Zipfian distribution.

few single machine stores running independently on the servers, in which case, the client needs to divide the key space (also called sharding) and send queries to the appropriate server. These experiments are denoted “client sharded” in the figures. The single machine key-value stores include a delegation based hashmap called trust here, Dashmap from Rust crate “dashmap” version “6.1.0” [1] and a hashmap with 4096 shards protected by a readers-writers lock referred as Rwlock below. In case of Redis, we disable replication and disk backup and provide two I/O threads per instance. Each server machine runs eight Redis instances. This configuration was selected after extensive experimentation to ensure maximum performance. The key space in these experiments consists of one million keys and queries consist of 8 byte keys and 16 byte values. Each reported data point here is an average of three runs after an initial warm up stage that populates all the keys. For these experiments up to 8 machines act as servers and up to 4 machines act as clients with each client machine running 256 threads

over 64 cores. For Hydra, each client thread connects to one server using TCP. However, for trust, Dashmap and Rwlock, each thread needs to connect with each server.

Figure 27 and Figure 28 show the number of operations per second for a read only workload and 5% write, 95% read workload respectively with a uniform distribution of keys. Hydra scales much better than running independent servers and leaving the sharding to the clients. This is explained by the fact that for Hydra, a client can send a whole batch of requests to any of the servers regardless of where the keys may reside, but for single machine stores, the client needs to divide the requests based on the keys and send them to appropriate servers. This reduces the batch size and increases the number of packets being sent over the network. That in turn results in the servers spending more time processing the TCP communication. This is evidenced by perf, the Linux performance analysis tool [40], showing that a trust server, which is also delegation based, spends 30% more time in kernel mode than a Hydra server. Dashmap and Rwlock do better than trust due to low contention but suffer from higher system call overhead much the same. Redis is a production ready key-value store that typically offers 400,000 to 600,000 operations per second depending on workload [39]. This is significantly less than per thread performance of all of the stores in this experiment. We also include a line labeled “redis hypothetical”, showing how it would scale were it to have the same per thread performance as Hydra.

Figure 29 and Figure 30 show the number of operations per second for a read only workload and 5% write, 95% read workload respectively with a Zipfian distribution of keys. We see similar trends for read only workload but Dashmap and Rwlock scale particularly poorly with the mixed workload due to high contention resulting from most of the queries being about a small number of keys. The contention stems from Rwlock and Dashmap, which

also uses an Rwlock internally, not allowing any readers while the value for a particular key is being updated. This is not an issue with either of the delegation based systems as reads and writes are treated the same way.

#### 5.4 Application: Single Source Shortest Path (SSSP)

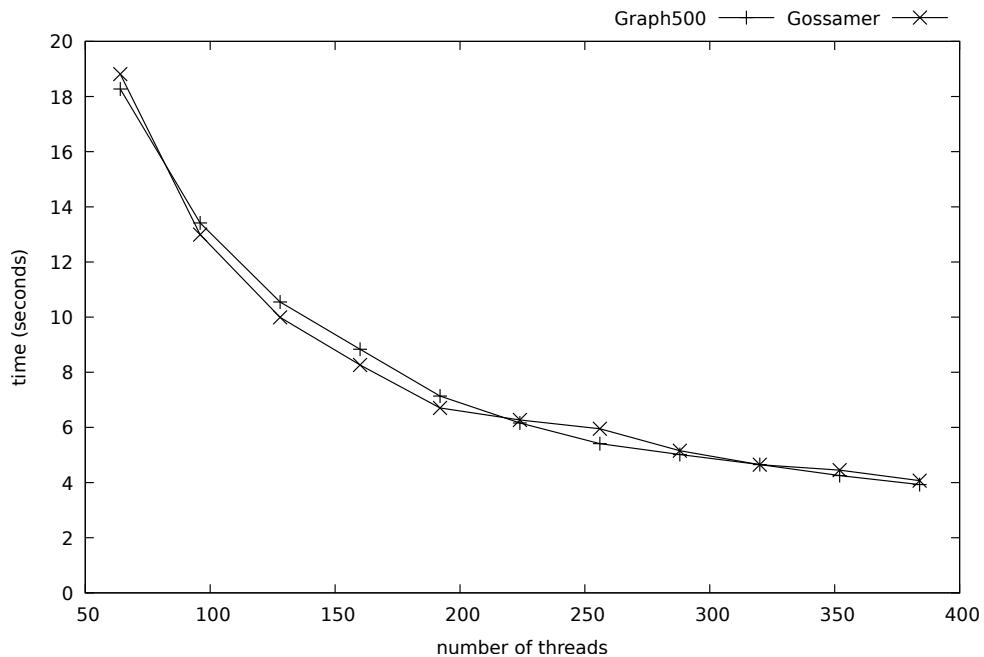


FIGURE 31. Time taken for SSSP on a graph with  $2^{27}$  nodes and 32 average degree.

For this benchmark, we implement a distributed version of Dijkstra’s algorithm for shortest path to each node in a weighted undirected graph given a source node. We compare this application to the graph500 reference benchmark [3] for supercomputers. Even though the graph500 benchmark uses the delta-stepping algorithm, we observed that both systems traverse roughly equal number of edges, making any algorithmic differences negligible. The graphs for these experiments were generated by graph500’s implementation of kronecker generator with an edge factor of 32. For each experiment we compute Single Source Shortest

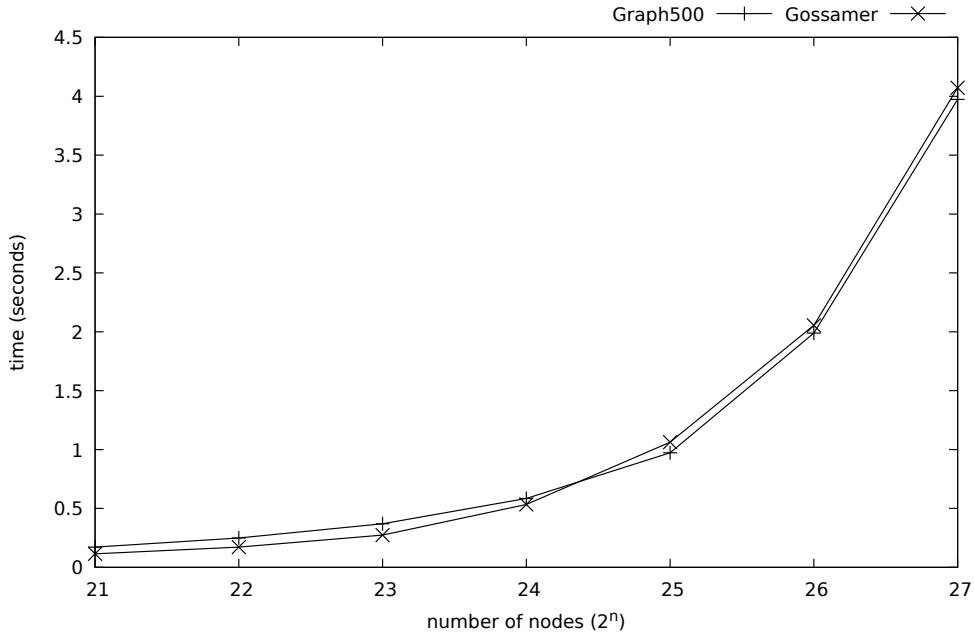


FIGURE 32. Time taken for SSSP on graphs with  $2^n$  nodes and 32 average degree.

Path (SSSP) on eight randomly chosen source nodes and report the average time taken per run. SSSP involves computing the shortest path to all reachable nodes in the graph, given a source node.

Figure 31 shows the results of SSSP on a graph with  $2^{27}$  nodes and average degree of 32, as the size of the cluster increases. Each successive data point increments the thread count by 32. Even though graph500 is a highly specialized application designed to benchmark supercomputers, *Gossamer* is still able to match its performance while remaining a general-purpose framework. Similarly, Figure 32 shows *Gossamer* achieving comparable performance to graph500 across a range of graph sizes with 384 threads across 12 machines.

## 6 Conclusion

In this chapter, we present *Gossamer*, a delegation based distributed programming model. *Gossamer* allows single machine applications using  $Trust\langle T \rangle$  to be distributed across a small rack with minimal changes to code. We take advantage of the fact that the same binary is run on all machines for a *Gossamer* process to reduce the request size and lower network overhead. We also demonstrate that *Gossamer* scales well. It can match the performance of a specialized benchmark despite being a general programming model, and it outperforms client sharded key value stores.

## CHAPTER 5

# Conclusion and Future Work

### 1 Conclusion

This dissertation highlights delegation as a viable alternative to locks for shared memory synchronization while also being applicable in a distributed setting. We introduced two system, *Trust*< $T$ > and *Gossamer*, to demonstrate the versatility of this approach across both domains. *Trust*< $T$ > is a multithreaded programming framework that can outperform traditional locks with high contention workloads. *Gossamer* is a distributed programming framework that scales well with increasing number of machines. Together, they illustrate how delegation unifies concurrency and distribution under a safe, efficient, and scalable programming paradigm, enabling high-performance computation without the complexities of shared memory or explicit synchronization.

### 2 Future Work

Below we discuss some avenues for further research and improvements to *Trust*< $T$ > and *Gossamer*.

#### 2.1 Better Profiling Tools

Due to the polling nature of trustees, it can be difficult to gauge how much useful work any given trustee is doing at a particular point in the application’s run time. This mismatch can cause expected performance to differ from observed performance. As an example, even

if some of the trustees are not getting any work, they will still show up as busy on profiling tools like htop and perf. Tools that can distinguish between idle spinning and productive work would greatly aid in diagnosing performance issues.

## 2.2 Better Load Balancing

Both *Trust*< $T$ > and *Gossamer* leave the responsibility for distributing data across trustees to the application developer. While this approach is workable for a proof-of-concept, it complicates application design: any change in workload distribution may require substantial data redistribution. This static assignment can also leave performance unused—for example, when a few trustees are heavily loaded while others remain mostly idle. A dynamic mechanism that detects imbalance at runtime and migrates entrusted data among trustees could potentially improve throughput and responsiveness. However, such a solution introduces diagnostic and migration overhead, which may be unacceptable when most or all trustees are already heavily loaded. Determining whether dynamic redistribution is beneficial, and under what conditions, requires further empirical investigation.

## 2.3 Handling Network Conditions

*Gossamer* relies on ideal network conditions—for example, environments with no congestion, symmetric routing, and no dropped packets. However, such perfectly reliable networks are rare in real-world deployments, which could significantly hinder the practical adoption of the protocol. To improve its robustness, additional mechanisms for managing adverse network conditions need to be explored. Potential strategies include implementing rate-limiting techniques to prevent overload, incorporating retransmission methods to recover from lost packets, and leveraging Explicit Congestion Notification (ECN) to proactively signal and

## Chapter 5. Conclusion and Future Work

respond to network congestion. Investigating these approaches presents valuable directions for future research and development.

## Bibliography

- [1] Dashmap. <https://crates.io/crates/dashmap/6.1.0>.
- [2] erpc throughput with 100 gbps connection. <https://github.com/erpc-io/eRPC/issues/32#issuecomment-510968083>.
- [3] Graph500 reference benchmark. <https://graph500.org>.
- [4] Redis cluster. [https://redis.io/docs/latest/operate/oss\\_and\\_stack/management/scaling/](https://redis.io/docs/latest/operate/oss_and_stack/management/scaling/).
- [5] A. Adamson and N. Williams. Remote procedure call (RPC) security version 3. Technical report, November 2016.
- [6] Yehuda Afek, Amir Levy, and Adam Morrison. Programming with hardware lock elision. *SIGPLAN Not.*, 48(8):295–296, February 2013.
- [7] Noaman Ahmad, Ben Baenen, Chen Chen, and Jakob Eriksson. Delegation with Trust<T>: A scalable, type- and memory-safe alternative to locks. *arXiv*, 2024.
- [8] Dan Alistarh, Keren Censor-Hillel, and Nir Shavit. Are lock-free concurrent algorithms practically wait-free? *Journal of the Association for Computing Machinery*, 63(4), 9 2016.
- [9] Hagit Attiya, Armando Castañeda, Danny Hendler, and Matthieu Perrin. Separating lock-freedom from wait-freedom at every level of the consensus hierarchy. *Journal of Parallel and Distributed Computing*, 163:181–197, 2022.
- [10] Pavan Balaji, Hemal Shah, and D.K. Panda. Sockets vs rdma interface over 10gigabit networks: An in-depth analysis of the memory traffic bottleneck. 01 2004.
- [11] Dotan Barak. Rdmmamojo. [https://www.rdmamojo.com/2013/02/15/ibv\\_poll\\_cq/](https://www.rdmamojo.com/2013/02/15/ibv_poll_cq/).
- [12] B. Bershad, T. Anderson, E. Lazowska, and H. Levy. Lightweight remote procedure call. In *Proceedings of the Twelfth ACM Symposium on Operating Systems Principles*, SOSP ’89, pages 102–113, New York, NY, USA, 1989. ACM.

## Chapter 5. Bibliography

- [13] Andrew D. Birrell and Bruce Jay Nelson. Implementing remote procedure calls. *ACM Trans. Comput. Syst.*, 2(1):39–59, February 1984.
- [14] Roy Frank Brabson, John J Majikes, and Janet C Wolf. Method and system for improved computer network efficiency in use of remote procedure call applications, March 22 2011. US Patent 7,913,262.
- [15] R. Brightwell and K. D. Underwood. An analysis of nic resource usage for offloading mpi. In *18th International Parallel and Distributed Processing Symposium, 2004. Proceedings.*, pages 183–, April 2004.
- [16] Ron Brightwell. A new mpi implementation for cray shmem. In Dieter Kranzlmüller, Péter Kacsuk, and Jack Dongarra, editors, *Recent Advances in Parallel Virtual Machine and Message Passing Interface*, pages 122–130, Berlin, Heidelberg, 2004. Springer Berlin Heidelberg.
- [17] Ron Brightwell, Rolf Riesen, and Keith D. Underwood. Analyzing the impact of overlap, offload, and independent progress for message passing interface applications. *The International Journal of High Performance Computing Applications*, 19(2):103–117, 2005.
- [18] Irina Calciu, Dave Dice, Tim Harris, Maurice Herlihy, Alex Kogan, Virendra Marathe, and Mark Moir. Message passing or shared memory: Evaluating the delegation abstraction for multicores. In *International Conference on Principles of Distributed Systems*, pages 83–97. Springer, 2013.
- [19] Chen Chen. *Optimizing Key-Value Store Systems Across Multiple Dimensions*. Phd thesis, University of Illinois Chicago, 2025.
- [20] M. Cohen, T. Ponte, S. Rossetto, and N. Rodriguez. Using coroutines for rpc in sensor networks. In *2007 IEEE International Parallel and Distributed Processing Symposium*, pages 1–8, March 2007.
- [21] Tudor David, Rachid Guerraoui, and Vasileios Trigonakis. Everything you always wanted to know about synchronization but were afraid to ask. In *Proceedings of the Twenty-Fourth ACM Symposium on Operating Systems Principles*, pages 33–48. ACM, 2013.
- [22] Dave Dice, Virendra J Marathe, and Nir Shavit. Flat-combining numa locks. In *Proceedings of the twenty-third annual ACM symposium on Parallelism in algorithms and architectures*, pages 65–74. ACM, 2011.

## Chapter 5. Bibliography

- [23] P.S. DODDS, D.H. ROTHMAN, and J.S. WEITZ. Re-examination of the “3/4-law” of metabolism. *Journal of Theoretical Biology*, 209(1):9–27, 2001.
- [24] Daniel C Doolan, Sabin Tabirca, and Laurence T Yang. Mmpi a message passing interface for the mobile environment. In *Proceedings of the 6th International Conference on Advances in Mobile Computing and Multimedia*, MoMM ’08, pages 317–321, New York, NY, USA, 2008. ACM.
- [25] Aleksandar Dragojević, Dushyanth Narayanan, Miguel Castro, and Orion Hodson. Farm: Fast remote memory. In *11th USENIX Symposium on Networked Systems Design and Implementation (NSDI 14)*, pages 401–414, Seattle, WA, 2014. USENIX Association.
- [26] Panagiota Fatourou and Nikolaos D Kallimanis. A highly-efficient wait-free universal construction. In *Proceedings of the twenty-third annual ACM symposium on Parallelism in algorithms and architectures*, pages 325–334. ACM, 2011.
- [27] Panagiota Fatourou and Nikolaos D Kallimanis. Revisiting the combining synchronization technique. In *ACM SIGPLAN Notices*, volume 47, pages 257–266. ACM, 2012.
- [28] Panagiota Fatourou and Eric Ruppert. Lock-Free Augmented Trees. In Dan Alistarh, editor, *38th International Symposium on Distributed Computing (DISC 2024)*, volume 319 of *Leibniz International Proceedings in Informatics (LIPIcs)*, pages 23:1–23:24, Dagstuhl, Germany, 2024. Schloss Dagstuhl – Leibniz-Zentrum für Informatik.
- [29] Message P Forum. Mpi: A message-passing interface standard. Technical report, USA, 1994.
- [30] Chryssis Georgiou, Nicolas Nicolaou, and Andria Trigeorgi. Dynamic distributed shared memory: When&how? In *Proceedings of the 12th Workshop on Principles and Practice of Consistency for Distributed Data*, PaPoC ’25, pages 1–7, New York, NY, USA, 2025. Association for Computing Machinery.
- [31] Vishal Gupta, Kumar Kartikeya Dwivedi, Yugesh Kothari, Yueyang Pan, Diyu Zhou, and Sanidhya Kashyap. Ship your critical section, not your data: Enabling transparent delegation with TCLOCKS. In *17th USENIX Symposium on Operating Systems Design and Implementation (OSDI 23)*, pages 1–16, Boston, MA, July 2023. USENIX Association.

## Chapter 5. Bibliography

- [32] Danny Hendler, Itai Incze, Nir Shavit, and Moran Tzafrir. Flat combining and the synchronization-parallelism tradeoff. In *Proceedings of the twenty-second annual ACM symposium on Parallelism in algorithms and architectures*, pages 355–364. ACM, 2010.
- [33] Vivien Quéma Hugo Guiroux, Renaud Lachaize. Multicore locks: The case is not closed yet. In *2016 USENIX Annual Technical Conference (USENIX ATC 16)*, pages 649–662, Denver, CO, June 2016. USENIX Association.
- [34] Jerry Neumann. Power Laws in Venture. <https://reactionwheel.net/2015/06/power-laws-in-venture.html>.
- [35] Anuj Kalia, Michael Kaminsky, and David Andersen. Datacenter RPCs can be general and fast. In *16th USENIX Symposium on Networked Systems Design and Implementation (NSDI 19)*, pages 1–16, Boston, MA, February 2019. USENIX Association.
- [36] Sanidhya Kashyap, Changwoo Min, and Taesoo Kim. Scalable NUMA-aware blocking synchronization primitives. In *2017 USENIX Annual Technical Conference (USENIX ATC 17)*, pages 603–615, Santa Clara, CA, July 2017. USENIX Association.
- [37] Alex Kogan and Erez Petrank. A methodology for creating fast wait-free data structures. *SIGPLAN Not.*, 47(8):141–150, February 2012.
- [38] Alex Kogan and Erez Petrank. A methodology for creating fast wait-free data structures. In *Proceedings of the 17th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming, PPoPP ’12*, page 141–150, New York, NY, USA, 2012. Association for Computing Machinery.
- [39] kostas. How many requests per second can i get out of redis? [https://skipperkongen.dk/2013/08/27/how-many-requests-per-second-can-i-get-out-of-redis/?utm\\_source=chatgpt.com](https://skipperkongen.dk/2013/08/27/how-many-requests-per-second-can-i-get-out-of-redis/?utm_source=chatgpt.com), 8 2013.
- [40] Linux Kernel Organization. perf: Linux profiling with performance counters. <https://man7.org/linux/man-pages/man1/perf.1.html>.
- [41] Luke Logan, Anthony Kougkas, and Xian-He Sun. Megammap: Blurring the boundary between memory and storage for data-intensive workloads. In *Proceedings of the International Conference for High Performance Computing, Networking, Storage, and Analysis, SC ’24*. IEEE Press, 2024.

- [42] Jean-Pierre Lozi, Florian David, Gaël Thomas, Julia L Lawall, Gilles Muller, et al. Remote core locking: Migrating critical-section execution to improve the performance of multithreaded applications. In *2012 USENIX Annual Technical Conference (USENIX ATC 12)*, pages 65–76, 2012.
- [43] Haoran Ma, Yifan Qiao, Shi Liu, Shan Yu, Yuanjiang Ni, Qingda Lu, Jiesheng Wu, Yiying Zhang, Miryung Kim, and Harry Xu. DRust: Language-Guided distributed shared memory with fine granularity, full transparency, and ultra efficiency. In *18th USENIX Symposium on Operating Systems Design and Implementation (OSDI 24)*, pages 97–115, Santa Clara, CA, July 2024. USENIX Association.
- [44] Michael Marty, Marc de Kruijf, Jacob Adriaens, Christopher Alfeld, Sean Bauer, Carlo Contavalli, Michael Dalton, Nandita Dukkipati, William C. Evans, Steve Gribble, and et al. Snap: A microkernel approach to host networking. In *Proceedings of the 27th ACM Symposium on Operating Systems Principles*, SOSP ’19, pages 399–413, New York, NY, USA, 2019. Association for Computing Machinery.
- [45] Phillip Merrick, Stewart O Allen, et al. Xml remote procedure call (xml-rpc), August 25 2015. US Patent 9,116,762.
- [46] Christopher Mitchell, Yifeng Geng, and Jinyang Li. Using one-sided RDMA reads to build a fast, cpu-efficient key-value store. In *Presented as part of the 2013 USENIX Annual Technical Conference (USENIX ATC 13)*, pages 103–114, San Jose, CA, 2013. USENIX.
- [47] Ruslan Nikolaev and Binoy Ravindran. Universal wait-free memory reclamation. In *Proceedings of the 25th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*, PPoPP ’20, page 130–143, New York, NY, USA, 2020. Association for Computing Machinery.
- [48] Ruslan Nikolaev and Binoy Ravindran. wcq: a fast wait-free queue with bounded memory usage. In *Proceedings of the 27th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*, PPoPP ’22, pages 461–462, New York, NY, USA, 2022. Association for Computing Machinery.
- [49] Michael I Ojovan and William E Lee. Topologically disordered systems at the glass transition. *Journal of Physics: Condensed Matter*, 18(50):11507, nov 2006.

- [50] Yoshihiro Oyama, Kenjiro Taura, and Akinori Yonezawa. Executing parallel programs with synchronization bottlenecks efficiently. In *Proceedings of the International Workshop on Parallel and Distributed Computing for Symbolic and Irregular Applications*, volume 16. Citeseer, 1999.
- [51] Xing Pan and Frank Mueller. Numa-aware memory coloring for multicore real-time systems. *Journal of Systems Architecture*, 118:102188, 2021.
- [52] A. Patil, A. Shah, S. Gaikwad, A. A. Mishra, S. S. Kohli, and S. Dhage. Fault tolerance in cluster computing system. In *2011 International Conference on P2P, Parallel, Grid, Cloud and Internet Computing*, pages 408–412, Oct 2011.
- [53] C. Peterson, J. Sutton, and P. Wiley. iwarpc: a 100-mops, liw microprocessor for multicomputers. *IEEE Micro*, 11(3):26–29, June 1991.
- [54] F. Petrini, , A. Hoisie, S. Coll, and E. Frachtenberg. The quadrics network (qsnet): high-performance clustering technology. In *HOT 9 Interconnects. Symposium on High Performance Interconnects*, pages 125–130, Aug 2001.
- [55] Darko Petrović, Thomas Ropars, and André Schiper. On the performance of delegation over cache-coherent shared memory. In *Proceedings of the 2015 International Conference on Distributed Computing and Networking*, page 17. ACM, 2015.
- [56] Gregory F Pfister. An introduction to the infiniband architecture. *High Performance Mass Storage and Parallel I/O*, 42:617–632, 2001.
- [57] Sepideh Roghanchi, Jakob Eriksson, and Nilanjana Basu. Ffwd: Delegation is (much) faster than you think. In *Proceedings of the 26th Symposium on Operating Systems Principles*, SOSP ’17, pages 342–358, New York, NY, USA, 2017. ACM.
- [58] Ori Shalev and Nir Shavit. Predictive log-synchronization. In *ACM SIGOPS Operating Systems Review*, volume 40, pages 305–315. ACM, 2006.
- [59] Nagaraj Shyam, Craig Harmer, and Ken Beck. Managing remote procedure calls when a server is unavailable, September 22 2015. US Patent 9,141,449.

## Chapter 5. Bibliography

- [60] J. Soumagne, D. Kimpe, J. Zounmevo, M. Chaarawi, Q. Koziol, A. Afsahi, and R. Ross. Mercury: Enabling remote procedure call for high-performance computing. In *2013 IEEE International Conference on Cluster Computing (CLUSTER)*, pages 1–8, Sep. 2013.
- [61] Jérôme Soumagne, D. Kimpe, Judicael A. Zounmevo, Mohamad Chaarawi, Q. Koziol, A. Afsahi, and R. Ross. Mercury: Enabling remote procedure call for high-performance computing. *2013 IEEE International Conference on Cluster Computing (CLUSTER)*, pages 1–8, 2013.
- [62] R. Srinivasan. Rpc: Remote procedure call protocol specification version 2, 1995.
- [63] Josef Stefan and Ludwig Boltzmann. On the relation between the emission and absorption of radiant heat. *Annalen der Physik und Chemie*, page 391–428, 1879.
- [64] H. Subramoni, P. Lai, M. Luo, and D. K. Panda. Rdma over ethernet — a preliminary study. In *2009 IEEE International Conference on Cluster Computing and Workshops*, pages 1–9, Aug 2009.
- [65] T. Talpey and B. Callaghan. Remote direct memory access transport for remote procedure call. Technical report, January 2010.
- [66] Markus Velten, Robert Schöne, Thomas Ilsche, and Daniel Hackenberg. Memory performance of amd epyc rome and intel cascade lake sp server processors. In *Proceedings of the 2022 ACM/SPEC on International Conference on Performance Engineering, ICPE ’22*, pages 165–175. ACM, April 2022.
- [67] Qing Wang, Youyou Lu, Erci Xu, Junru Li, Youmin Chen, and Jiwu Shu. Concordia: Distributed shared memory with In-Network cache coherence. In *19th USENIX Conference on File and Storage Technologies (FAST 21)*, pages 277–292. USENIX Association, February 2021.
- [68] Ruihong Wang, Jianguo Wang, Stratos Idreos, M. Tamer Özsu, and Walid G. Aref. The case for distributed shared-memory databases with rdma-enabled memory disaggregation. *Proc. VLDB Endow.*, 16(1):15–22, September 2022.
- [69] xacrimon. dashmap. <https://crates.io/crates/dashmap>. Version 5.5.3.
- [70] Hanmei Yang, Xin Zhao, Jin Zhou, Wei Wang, Sandip Kundu, Bo Wu, Hui Guan, and Tongping Liu. Numalloc: A faster numa memory allocator. In *Proceedings of the 2023 ACM SIGPLAN International*

## Chapter 5. Bibliography

- Symposium on Memory Management*, ISMM 2023, page 97–110, New York, NY, USA, 2023. Association for Computing Machinery.
- [71] Pen-Chung Yew, Nian-Feng Tzeng, et al. Distributing hot-spot addressing in large-scale multiprocessors. *IEEE Transactions on Computers*, 100:388–395, 1987.
- [72] Mingxing Zhang, Teng Ma, Jinqi Hua, Zheng Liu, Kang Chen, Ning Ding, Fan Du, Jinlei Jiang, Tao Ma, and Yongwei Wu. Partial failure resilient memory management system for (cxl-based) distributed shared memory. In *Proceedings of the 29th Symposium on Operating Systems Principles*, SOSP ’23, pages 658–674, New York, NY, USA, 2023. Association for Computing Machinery.
- [73] George K. Zipf. *Human Behavior and the Principle of Least Effort: An Introduction to Human Ecology*. Addison-Wesley, 1949.

## VITA

<b>NAME</b>	Noaman Ahmad
<b>EDUCATION</b>	BS, Computer Science, Lahore University of Management Sciences, Lahore, Pakistan, 2018 MS, Computer Science, University of Illinois Chicago, Chicago, Illinois, 2025
<b>TEACHING</b>	Hands-on Rust: A Practical Introduction (CS 194, Summer 2025), Primary Instructor Machine Organization (CS 261, Spring 2021), TA Operating Systems Design and Implementation (CS 461, Fall 2024 & Fall 2025), TA Systems Programming (CS 361, Spring 2025), TA
<b>PUBLICATIONS</b>	Noaman Ahmad, Ben Baenen, Chen Chen and Jakob Eriksson “Delegation with Trust<T>: A Scalable, Type-and Memory-Safe Alternative to Locks” In arXiv preprint 2408.11173 (2024). Noaman Ahmad, Ben Baenen, Chen Chen and Jakob Eriksson “Delegation with Trust” In preparation. Noaman Ahmad, Chen Chen and Jakob Eriksson “Gossamer: A novel programming paradigm for rack-wide applications” In preparation.